

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
import re
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
```

Madrid 2001

In [2]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\madrid_2001.csv")
a
```

Out[2]:

| | date | BEN | CO | EBE | MXV | NMHC | NO_2 | NOx | OXY | O_3 | |
|--------|---------------------|-------|------|------|-------|------|-----------|------------|------|-----------|-----|
| 0 | 2001-08-01 01:00:00 | NaN | 0.37 | NaN | NaN | NaN | 58.400002 | 87.150002 | NaN | 34.529999 | 10 |
| 1 | 2001-08-01 01:00:00 | 1.50 | 0.34 | 1.49 | 4.10 | 0.07 | 56.250000 | 75.169998 | 2.11 | 42.160000 | 10 |
| 2 | 2001-08-01 01:00:00 | NaN | 0.28 | NaN | NaN | NaN | 50.660000 | 61.380001 | NaN | 46.310001 | 10 |
| 3 | 2001-08-01 01:00:00 | NaN | 0.47 | NaN | NaN | NaN | 69.790001 | 73.449997 | NaN | 40.650002 | 6 |
| 4 | 2001-08-01 01:00:00 | NaN | 0.39 | NaN | NaN | NaN | 22.830000 | 24.799999 | NaN | 66.309998 | 7 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 217867 | 2001-04-01 00:00:00 | 10.45 | 1.81 | NaN | NaN | NaN | 73.000000 | 264.399994 | NaN | 5.200000 | 4 |
| 217868 | 2001-04-01 00:00:00 | 5.20 | 0.69 | 4.56 | NaN | 0.13 | 71.080002 | 129.300003 | NaN | 13.460000 | 2 |
| 217869 | 2001-04-01 00:00:00 | 0.49 | 1.09 | NaN | 1.00 | 0.19 | 76.279999 | 128.399994 | 0.35 | 5.020000 | 4 |
| 217870 | 2001-04-01 00:00:00 | 5.62 | 1.01 | 5.04 | 11.38 | NaN | 80.019997 | 197.000000 | 2.58 | 5.840000 | 3 |
| 217871 | 2001-04-01 00:00:00 | 8.09 | 1.62 | 6.66 | 13.04 | 0.18 | 76.809998 | 206.300003 | 5.20 | 8.340000 | 3 |

217872 rows × 16 columns



In [3]:

`a.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 217872 entries, 0 to 217871
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        217872 non-null object  
 1   BEN         70389 non-null  float64
 2   CO          216341 non-null float64
 3   EBE         57752 non-null  float64
 4   MXY         42753 non-null  float64
 5   NMHC        85719 non-null  float64
 6   NO_2        216331 non-null float64
 7   NOx         216318 non-null float64
 8   OXY         42856 non-null  float64
 9   O_3         216514 non-null float64
10  PM10        207776 non-null float64
11  PXY         42845 non-null  float64
12  SO_2        216403 non-null float64
13  TCH         85797 non-null  float64
14  TOL         70196 non-null  float64
15  station     217872 non-null int64  
dtypes: float64(14), int64(1), object(1)
memory usage: 26.6+ MB
```

In [4]:

```
b=a.fillna(value=87)
b
```

Out[4]:

| | date | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 |
|--------|---------------------|-------|------|-------|-------|-------|-----------|------------|-------|-----------|
| 0 | 2001-08-01 01:00:00 | 87.00 | 0.37 | 87.00 | 87.00 | 87.00 | 58.400002 | 87.150002 | 87.00 | 34.529999 |
| 1 | 2001-08-01 01:00:00 | 1.50 | 0.34 | 1.49 | 4.10 | 0.07 | 56.250000 | 75.169998 | 2.11 | 42.160000 |
| 2 | 2001-08-01 01:00:00 | 87.00 | 0.28 | 87.00 | 87.00 | 87.00 | 50.660000 | 61.380001 | 87.00 | 46.310001 |
| 3 | 2001-08-01 01:00:00 | 87.00 | 0.47 | 87.00 | 87.00 | 87.00 | 69.790001 | 73.449997 | 87.00 | 40.650002 |
| 4 | 2001-08-01 01:00:00 | 87.00 | 0.39 | 87.00 | 87.00 | 87.00 | 22.830000 | 24.799999 | 87.00 | 66.309998 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 217867 | 2001-04-01 00:00:00 | 10.45 | 1.81 | 87.00 | 87.00 | 87.00 | 73.000000 | 264.399994 | 87.00 | 5.200000 |
| 217868 | 2001-04-01 00:00:00 | 5.20 | 0.69 | 4.56 | 87.00 | 0.13 | 71.080002 | 129.300003 | 87.00 | 13.460000 |
| 217869 | 2001-04-01 00:00:00 | 0.49 | 1.09 | 87.00 | 1.00 | 0.19 | 76.279999 | 128.399994 | 0.35 | 5.020000 |
| 217870 | 2001-04-01 00:00:00 | 5.62 | 1.01 | 5.04 | 11.38 | 87.00 | 80.019997 | 197.000000 | 2.58 | 5.840000 |
| 217871 | 2001-04-01 00:00:00 | 8.09 | 1.62 | 6.66 | 13.04 | 0.18 | 76.809998 | 206.300003 | 5.20 | 8.340000 |

217872 rows × 16 columns

In [5]:

```
b.columns
```

Out[5]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
      'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

In [6]:

```
c=b.head(10)
c
```

Out[6]:

| | date | BEN | CO | EBE | MXV | NMHC | NO_2 | NOx | OXY | O_3 | I |
|---|---------------------|-------|------|-------|-------|-------|-----------|------------|-------|-----------|--------|
| 0 | 2001-08-01 01:00:00 | 87.00 | 0.37 | 87.00 | 87.00 | 87.00 | 58.400002 | 87.150002 | 87.00 | 34.529999 | 105.00 |
| 1 | 2001-08-01 01:00:00 | 1.50 | 0.34 | 1.49 | 4.10 | 0.07 | 56.250000 | 75.169998 | 2.11 | 42.160000 | 100.50 |
| 2 | 2001-08-01 01:00:00 | 87.00 | 0.28 | 87.00 | 87.00 | 87.00 | 50.660000 | 61.380001 | 87.00 | 46.310001 | 100.00 |
| 3 | 2001-08-01 01:00:00 | 87.00 | 0.47 | 87.00 | 87.00 | 87.00 | 69.790001 | 73.449997 | 87.00 | 40.650002 | 69.77 |
| 4 | 2001-08-01 01:00:00 | 87.00 | 0.39 | 87.00 | 87.00 | 87.00 | 22.830000 | 24.799999 | 87.00 | 66.309998 | 75.10 |
| 5 | 2001-08-01 01:00:00 | 2.11 | 0.63 | 2.48 | 5.94 | 0.05 | 66.260002 | 118.099998 | 3.15 | 33.500000 | 122.60 |
| 6 | 2001-08-01 01:00:00 | 87.00 | 0.28 | 87.00 | 87.00 | 87.00 | 35.799999 | 39.590000 | 87.00 | 68.250000 | 124.90 |
| 7 | 2001-08-01 01:00:00 | 87.00 | 0.67 | 87.00 | 87.00 | 87.00 | 74.830002 | 112.000000 | 87.00 | 26.410000 | 113.00 |
| 8 | 2001-08-01 01:00:00 | 87.00 | 0.41 | 87.00 | 87.00 | 87.00 | 33.209999 | 37.299999 | 87.00 | 62.299999 | 125.30 |
| 9 | 2001-08-01 01:00:00 | 87.00 | 0.17 | 87.00 | 87.00 | 0.13 | 24.129999 | 36.970001 | 87.00 | 46.200001 | 95.50 |



In [7]:

```
d=c[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
    'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]  
d
```

Out[7]:


| | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 | PM10 | PXY | SO_2 | TCH | TOL | station |
|---|-------|------|-------|-------|-------|-----------|------------|-------|-----------|------------|-------|------|-------|-------|---------|
| 0 | 87.00 | 0.37 | 87.00 | 87.00 | 87.00 | 58.400002 | 87.150002 | 87.00 | 34.529999 | 105.000000 | 87.00 | 0.37 | 87.00 | 87.00 | 0 |
| 1 | 1.50 | 0.34 | 1.49 | 4.10 | 0.07 | 56.250000 | 75.169998 | 2.11 | 42.160000 | 100.599998 | 1.50 | 0.34 | 1.49 | 0.07 | 1 |
| 2 | 87.00 | 0.28 | 87.00 | 87.00 | 87.00 | 50.660000 | 61.380001 | 87.00 | 46.310001 | 100.099998 | 87.00 | 0.28 | 87.00 | 87.00 | 2 |
| 3 | 87.00 | 0.47 | 87.00 | 87.00 | 87.00 | 69.790001 | 73.449997 | 87.00 | 40.650002 | 69.779999 | 87.00 | 0.47 | 87.00 | 87.00 | 3 |
| 4 | 87.00 | 0.39 | 87.00 | 87.00 | 87.00 | 22.830000 | 24.799999 | 87.00 | 66.309998 | 75.180000 | 87.00 | 0.39 | 87.00 | 87.00 | 4 |
| 5 | 2.11 | 0.63 | 2.48 | 5.94 | 0.05 | 66.260002 | 118.099998 | 3.15 | 33.500000 | 122.699997 | 2.11 | 0.63 | 2.48 | 0.05 | 5 |
| 6 | 87.00 | 0.28 | 87.00 | 87.00 | 87.00 | 35.799999 | 39.590000 | 87.00 | 68.250000 | 124.900002 | 87.00 | 0.28 | 87.00 | 87.00 | 6 |
| 7 | 87.00 | 0.67 | 87.00 | 87.00 | 87.00 | 74.830002 | 112.000000 | 87.00 | 26.410000 | 113.000000 | 87.00 | 0.67 | 87.00 | 87.00 | 7 |
| 8 | 87.00 | 0.41 | 87.00 | 87.00 | 87.00 | 33.209999 | 37.299999 | 87.00 | 62.299999 | 125.300003 | 87.00 | 0.41 | 87.00 | 87.00 | 8 |
| 9 | 87.00 | 0.17 | 87.00 | 87.00 | 0.13 | 24.129999 | 36.970001 | 87.00 | 46.200001 | 95.589996 | 87.00 | 0.17 | 87.00 | 87.00 | 9 |

In [8]:

```
d.plot.line()
```

Out[8]:

<AxesSubplot:>

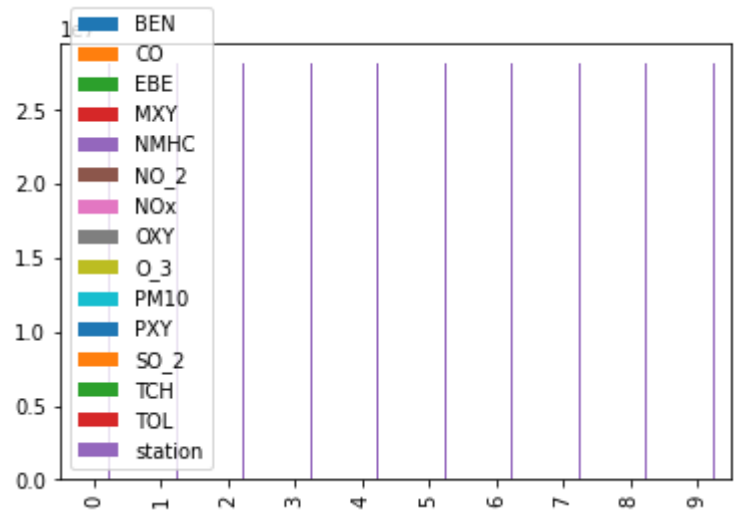


In [9]:

```
d.plot.bar()
```

Out[9]:

<AxesSubplot:>

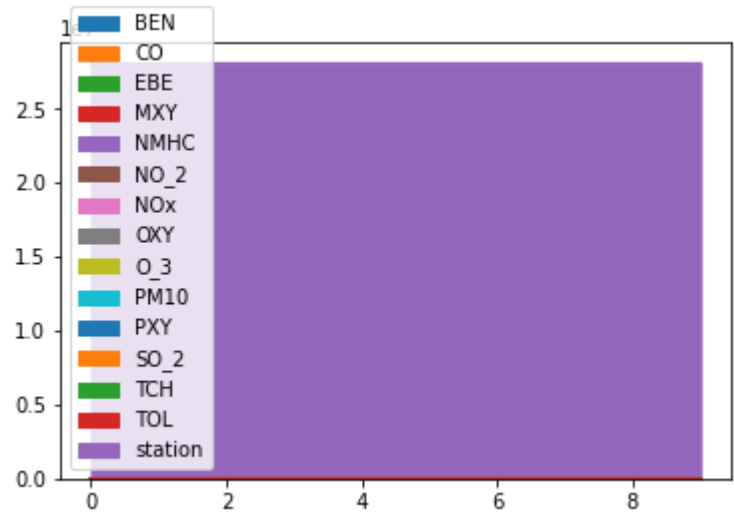


In [10]:

```
d.plot.area()
```

Out[10]:

<AxesSubplot:>

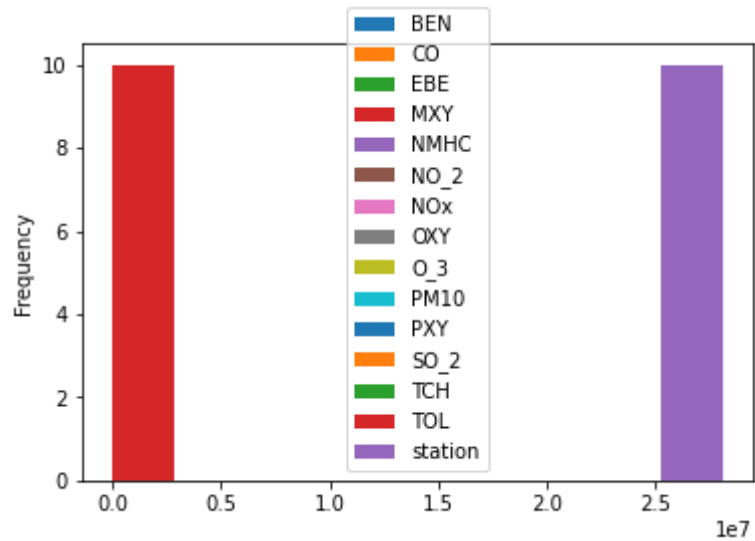


In [11]:

```
d.plot.hist()
```

Out[11]:

<AxesSubplot:ylabel='Frequency'>




In [12]:

```
d.plot.box()
```

Out[12]:

<AxesSubplot:>




In [13]:

```
d.plot.scatter(x='NO_2',y='NOx')
```

Out[13]:

<AxesSubplot:xlabel='NO_2', ylabel='NOx'>




In [14]:

```
d.plot.pie(y='NOx',figsize=(5,5))
```

Out[14]:

<AxesSubplot:ylabel='NOx'>




In [15]:

```
sns.pairplot(d)
```

Out[15]:

<seaborn.axisgrid.PairGrid at 0x1cbe59fb0d0>



In [16]:


```
sns.distplot(a['NOx'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:
FutureWarning: `distplot` is a deprecated function and will be removed in
a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level fun
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[16]:

```
<AxesSubplot:xlabel='NOx', ylabel='Density'>
```




In [17]:

```
sns.heatmap(d.corr())
```

Out[17]:

```
<AxesSubplot:>
```



In [18]:

```
x=d[ ['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY']]  
y=d[ 'TCH']
```

In [19]:

```
from sklearn.model_selection import train_test_split  
  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [20]:

```
from sklearn.linear_model import LinearRegression  
  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[20]:

LinearRegression()

In [21]:

```
print(lr.intercept_)  
  
1.311960462608127
```

In [22]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=[ 'Co-efficient' ])  
coeff
```

Out[22]:


| | Co-efficient |
|-------------|---------------|
| BEN | 0.000000e+00 |
| CO | -6.211505e-14 |
| EBE | 0.000000e+00 |
| MXY | 0.000000e+00 |
| NMHC | 9.849200e-01 |
| NO_2 | 7.333005e-16 |
| NOx | -9.633179e-17 |
| OXY | 0.000000e+00 |

In [23]:

```
prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[23]:

<matplotlib.collections.PathCollection at 0x1cbf3a3d130>



In [24]:

```
print(lr.score(x_test,y_test))
```

0.9999924406308133

In [25]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [26]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[26]:

Ridge(alpha=10)

In [27]:

```
rr.score(x_test,y_test)
```

Out[27]:

0.9999685317021586

In [28]:

```
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[28]:

Lasso(alpha=10)

In [29]:

```
la.score(x_test,y_test)
```

Out[29]:

0.9996339281315851

In [30]:

```
a1=b.head(7000)
a1
```

Out[30]:

| | date | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 |
|------|---------------------|-------|------|-----------|------|-------|-----------|-----------|-------|-----------|
| 0 | 2001-08-01 01:00:00 | 87.00 | 0.37 | 87.000000 | 87.0 | 87.00 | 58.400002 | 87.150002 | 87.00 | 34.529999 |
| 1 | 2001-08-01 01:00:00 | 1.50 | 0.34 | 1.490000 | 4.1 | 0.07 | 56.250000 | 75.169998 | 2.11 | 42.160000 |
| 2 | 2001-08-01 01:00:00 | 87.00 | 0.28 | 87.000000 | 87.0 | 87.00 | 50.660000 | 61.380001 | 87.00 | 46.310001 |
| 3 | 2001-08-01 01:00:00 | 87.00 | 0.47 | 87.000000 | 87.0 | 87.00 | 69.790001 | 73.449997 | 87.00 | 40.650002 |
| 4 | 2001-08-01 01:00:00 | 87.00 | 0.39 | 87.000000 | 87.0 | 87.00 | 22.830000 | 24.799999 | 87.00 | 66.309998 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 6995 | 2001-08-13 04:00:00 | 87.00 | 0.00 | 87.000000 | 87.0 | 0.08 | 18.580000 | 18.590000 | 87.00 | 56.660000 |
| 6996 | 2001-08-13 04:00:00 | 87.00 | 0.09 | 87.000000 | 87.0 | 87.00 | 29.580000 | 32.770000 | 87.00 | 52.709999 |
| 6997 | 2001-08-13 04:00:00 | 1.38 | 0.17 | 30.530001 | 87.0 | 0.25 | 54.880001 | 68.870003 | 87.00 | 23.240000 |
| 6998 | 2001-08-13 04:00:00 | 87.00 | 0.01 | 87.000000 | 87.0 | 87.00 | 19.580000 | 20.990000 | 87.00 | 51.270000 |
| 6999 | 2001-08-13 04:00:00 | 87.00 | 0.00 | 87.000000 | 87.0 | 0.05 | 17.200001 | 18.219999 | 87.00 | 38.090000 |

7000 rows × 16 columns



In [31]:

```
e=a1[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
      'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
e
```

Out[31]:

| | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 | PM1 |
|------|-------|------|-----------|------|-------|-----------|-----------|-------|-----------|-----------|
| 0 | 87.00 | 0.37 | 87.000000 | 87.0 | 87.00 | 58.400002 | 87.150002 | 87.00 | 34.529999 | 105.00000 |
| 1 | 1.50 | 0.34 | 1.490000 | 4.1 | 0.07 | 56.250000 | 75.169998 | 2.11 | 42.160000 | 100.59999 |
| 2 | 87.00 | 0.28 | 87.000000 | 87.0 | 87.00 | 50.660000 | 61.380001 | 87.00 | 46.310001 | 100.09999 |
| 3 | 87.00 | 0.47 | 87.000000 | 87.0 | 87.00 | 69.790001 | 73.449997 | 87.00 | 40.650002 | 69.77999 |
| 4 | 87.00 | 0.39 | 87.000000 | 87.0 | 87.00 | 22.830000 | 24.799999 | 87.00 | 66.309998 | 75.18000 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 6995 | 87.00 | 0.00 | 87.000000 | 87.0 | 0.08 | 18.580000 | 18.590000 | 87.00 | 56.660000 | 22.82000 |
| 6996 | 87.00 | 0.09 | 87.000000 | 87.0 | 87.00 | 29.580000 | 32.770000 | 87.00 | 52.709999 | 38.47000 |
| 6997 | 1.38 | 0.17 | 30.530001 | 87.0 | 0.25 | 54.880001 | 68.870003 | 87.00 | 23.240000 | 18.18000 |
| 6998 | 87.00 | 0.01 | 87.000000 | 87.0 | 87.00 | 19.580000 | 20.990000 | 87.00 | 51.270000 | 33.83000 |
| 6999 | 87.00 | 0.00 | 87.000000 | 87.0 | 0.05 | 17.200001 | 18.219999 | 87.00 | 38.090000 | 43.45999 |

7000 rows × 15 columns

In [32]:

```
f=e.iloc[:,0:14]
g=e.iloc[:, -1]
```

In [33]:

```
h=StandardScaler().fit_transform(f)
```

In [34]:

```
logr=LogisticRegression(max_iter=10000)
logr.fit(h,g)
```

Out[34]:

```
LogisticRegression(max_iter=10000)
```

In [49]:

```
from sklearn.model_selection import train_test_split
h_train,h_test,g_train,g_test=train_test_split(h,g,test_size=0.3)
```

In [50]:

```
i=[[10,20,30,40,50,60,11,22,33,44,55,54,21,78]]
```

In [51]:

```
prediction=logr.predict(i)
print(prediction)
```

```
[28079021]
```

In [52]:

```
logr.classes_
```

Out[52]:

```
array([28079001, 28079003, 28079004, 28079006, 28079007, 28079009,
       28079011, 28079012, 28079014, 28079015, 28079016, 28079018,
       28079019, 28079021, 28079022, 28079023, 28079024, 28079025,
       28079035, 28079036, 28079038, 28079039, 28079040, 28079099],
      dtype=int64)
```

In [53]:

```
logr.predict_proba(i)[0][0]
```

Out[53]:

```
2.528560047135413e-268
```

In [54]:

```
logr.predict_proba(i)[0][1]
```

Out[54]:

```
4.8254923316380694e-139
```

In [55]:

```
logr.score(h_test,g_test)
```

Out[55]:

```
0.6261904761904762
```

In [40]:

```
from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[40]:

```
ElasticNet()
```


In [41]:

```
print(en.coef_)
```

```
[0.      0.      0.      0.      0.98384654 0.  
 0.      0.      ]
```

In [42]:

```
print(en.intercept_)
```

```
1.3920293645489608
```

In [43]:

```
prediction=en.predict(x_test)  
print(en.score(x_test,y_test))
```

```
0.9999809149884327
```

In [56]:

```
from sklearn.ensemble import RandomForestClassifier  
  
rfc=RandomForestClassifier()  
rfc.fit(h_train,g_train)
```

Out[56]:

```
RandomForestClassifier()
```

In [57]:

```
parameters={'max_depth':[1,2,3,4,5],  
            'min_samples_leaf':[5,10,15,20,25],  
            'n_estimators':[10,20,30,40,50]  
            }
```

In [58]:

```
from sklearn.model_selection import GridSearchCV  
  
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(h_train,g_train)
```

Out[58]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
             param_grid={'max_depth': [1, 2, 3, 4, 5],  
                         'min_samples_leaf': [5, 10, 15, 20, 25],  
                         'n_estimators': [10, 20, 30, 40, 50]},  
             scoring='accuracy')
```

In [59]:

```
grid_search.best_score_
```

Out[59]:

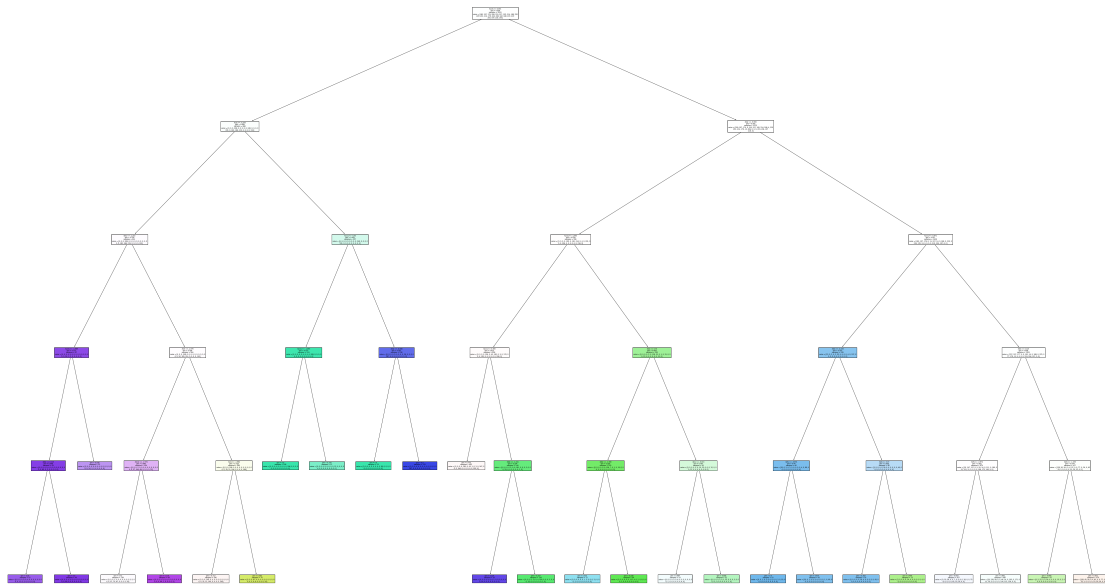
```
0.633061224489796
```

In [60]:

```
rfc_best=grid_search.best_estimator_
```

In [78]:

```
from sklearn.tree import plot_tree  
  
plt.figure(figsize=(80,50))  
plot_tree(rfc_best.estimators_[2],filled=True)
```



Conclusion: Linear score=0.9999924406308133, Ridge score=0.9999685317021586, Lasso Score=0.9996339281315851, Elastic Score=0.9999809149884327, Logistic Score=0.6261904761904762, RandomForest score=0.633061224489796.

Hence, Linear Regression model has the highest accuracy.

Madrid 2002

In [79]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\madrid_2002.csv")
a
```

Out[79]:

| | date | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 | P |
|--------|---------------------|------|------|------|-------|------|------------|------------|------|-------|--------|
| 0 | 2002-04-01 01:00:00 | NaN | 1.39 | NaN | NaN | NaN | 145.100006 | 352.100006 | NaN | 6.54 | 41.990 |
| 1 | 2002-04-01 01:00:00 | 1.93 | 0.71 | 2.33 | 6.20 | 0.15 | 98.150002 | 153.399994 | 2.67 | 6.85 | 20.980 |
| 2 | 2002-04-01 01:00:00 | NaN | 0.80 | NaN | NaN | NaN | 103.699997 | 134.000000 | NaN | 13.01 | 28.440 |
| 3 | 2002-04-01 01:00:00 | NaN | 1.61 | NaN | NaN | NaN | 97.599998 | 268.000000 | NaN | 5.12 | 42.180 |
| 4 | 2002-04-01 01:00:00 | NaN | 1.90 | NaN | NaN | NaN | 92.089996 | 237.199997 | NaN | 7.28 | 76.330 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 217291 | 2002-11-01 00:00:00 | 4.16 | 1.14 | NaN | NaN | NaN | 81.080002 | 265.700012 | NaN | 7.21 | 36.750 |
| 217292 | 2002-11-01 00:00:00 | 3.67 | 1.73 | 2.89 | NaN | 0.38 | 113.900002 | 373.100006 | NaN | 5.66 | 63.380 |
| 217293 | 2002-11-01 00:00:00 | 1.37 | 0.58 | 1.17 | 2.37 | 0.15 | 65.389999 | 107.699997 | 1.30 | 9.11 | 9.640 |
| 217294 | 2002-11-01 00:00:00 | 4.51 | 0.91 | 4.83 | 10.99 | NaN | 149.800003 | 202.199997 | 1.00 | 5.75 | 1 |
| 217295 | 2002-11-01 00:00:00 | 3.11 | 1.17 | 3.00 | 7.77 | 0.26 | 80.110001 | 180.300003 | 2.25 | 7.38 | 29.240 |

217296 rows × 16 columns



In [80]:

```
a=a.head(1000)
a
```

Out[80]:

| | date | BEN | CO | EBE | MXV | NMHC | NO_2 | NOx | OXY | O_3 | P |
|-----|---------------------|------|------|------|------|------|------------|------------|------|-----------|--------|
| 0 | 2002-04-01 01:00:00 | NaN | 1.39 | NaN | NaN | NaN | 145.100006 | 352.100006 | NaN | 6.540000 | 41.990 |
| 1 | 2002-04-01 01:00:00 | 1.93 | 0.71 | 2.33 | 6.20 | 0.15 | 98.150002 | 153.399994 | 2.67 | 6.850000 | 20.980 |
| 2 | 2002-04-01 01:00:00 | NaN | 0.80 | NaN | NaN | NaN | 103.699997 | 134.000000 | NaN | 13.010000 | 28.440 |
| 3 | 2002-04-01 01:00:00 | NaN | 1.61 | NaN | NaN | NaN | 97.599998 | 268.000000 | NaN | 5.120000 | 42.180 |
| 4 | 2002-04-01 01:00:00 | NaN | 1.90 | NaN | NaN | NaN | 92.089996 | 237.199997 | NaN | 7.280000 | 76.330 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 995 | 2002-04-02 16:00:00 | 2.19 | 0.36 | NaN | NaN | NaN | 58.709999 | 93.349998 | NaN | 70.830002 | 19.850 |
| 996 | 2002-04-02 16:00:00 | 0.87 | 0.30 | 1.00 | NaN | 0.09 | 32.580002 | 45.150002 | NaN | 77.910004 | 14.430 |
| 997 | 2002-04-02 16:00:00 | 0.46 | 0.50 | 0.27 | 0.41 | 0.10 | 11.550000 | 13.290000 | 0.49 | 82.260002 | 19.800 |
| 998 | 2002-04-02 16:00:00 | 1.11 | 0.44 | 0.96 | 2.00 | NaN | 96.790001 | 209.600006 | 0.66 | 34.540001 | |
| 999 | 2002-04-02 16:00:00 | 1.98 | 0.45 | 2.12 | 6.27 | 0.11 | 49.419998 | 75.730003 | 2.86 | 76.000000 | 25.160 |

1000 rows × 16 columns



In [81]:

```
b=a.dropna()  
b
```

Out[81]:

| | date | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 |
|-----|---------------------|------|------|------|-----------|------|------------|------------|------|-----------|
| 1 | 2002-04-01 01:00:00 | 1.93 | 0.71 | 2.33 | 6.200000 | 0.15 | 98.150002 | 153.399994 | 2.67 | 6.850000 |
| 5 | 2002-04-01 01:00:00 | 3.19 | 0.72 | 3.23 | 7.650000 | 0.11 | 113.699997 | 187.000000 | 3.53 | 12.370000 |
| 22 | 2002-04-01 01:00:00 | 2.02 | 0.80 | 1.57 | 3.660000 | 0.15 | 93.860001 | 101.300003 | 1.77 | 6.990000 |
| 24 | 2002-04-01 01:00:00 | 3.02 | 1.04 | 2.43 | 5.380000 | 0.21 | 103.699997 | 195.399994 | 2.15 | 14.040000 |
| 26 | 2002-04-01 02:00:00 | 2.02 | 0.53 | 2.24 | 5.970000 | 0.12 | 91.599998 | 136.199997 | 2.55 | 6.760000 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 974 | 2002-04-02 15:00:00 | 2.09 | 0.57 | 2.17 | 5.880000 | 0.15 | 51.470001 | 83.099998 | 2.59 | 71.410004 |
| 976 | 2002-04-02 16:00:00 | 1.96 | 0.40 | 2.09 | 5.440000 | 0.07 | 57.000000 | 83.410004 | 2.35 | 67.790001 |
| 980 | 2002-04-02 16:00:00 | 5.06 | 0.63 | 6.80 | 17.209999 | 0.09 | 69.510002 | 141.600006 | 7.95 | 67.720001 |
| 997 | 2002-04-02 16:00:00 | 0.46 | 0.50 | 0.27 | 0.410000 | 0.10 | 11.550000 | 13.290000 | 0.49 | 82.260002 |
| 999 | 2002-04-02 16:00:00 | 1.98 | 0.45 | 2.12 | 6.270000 | 0.11 | 49.419998 | 75.730003 | 2.86 | 76.000000 |

160 rows × 16 columns

In [82]:

```
b.columns
```

Out[82]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
      'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],  
      dtype='object')
```

In [83]:

```
b=b[['BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
      'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
b
```

Out[83]:

| | BEN | CO | EBE | MXV | NMHC | NO_2 | NOx | OXY | O_3 | PM10 |
|-----|------|------|------|-----------|------|------------|------------|------|-----------|-----------|
| 1 | 1.93 | 0.71 | 2.33 | 6.200000 | 0.15 | 98.150002 | 153.399994 | 2.67 | 6.850000 | 20.980000 |
| 5 | 3.19 | 0.72 | 3.23 | 7.650000 | 0.11 | 113.699997 | 187.000000 | 3.53 | 12.370000 | 27.450001 |
| 22 | 2.02 | 0.80 | 1.57 | 3.660000 | 0.15 | 93.860001 | 101.300003 | 1.77 | 6.990000 | 33.000000 |
| 24 | 3.02 | 1.04 | 2.43 | 5.380000 | 0.21 | 103.699997 | 195.399994 | 2.15 | 14.040000 | 37.310001 |
| 26 | 2.02 | 0.53 | 2.24 | 5.970000 | 0.12 | 91.599998 | 136.199997 | 2.55 | 6.760000 | 19.980000 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 974 | 2.09 | 0.57 | 2.17 | 5.880000 | 0.15 | 51.470001 | 83.099998 | 2.59 | 71.410004 | 31.059999 |
| 976 | 1.96 | 0.40 | 2.09 | 5.440000 | 0.07 | 57.000000 | 83.410004 | 2.35 | 67.790001 | 25.690001 |
| 980 | 5.06 | 0.63 | 6.80 | 17.209999 | 0.09 | 69.510002 | 141.600006 | 7.95 | 67.720001 | 21.950001 |
| 997 | 0.46 | 0.50 | 0.27 | 0.410000 | 0.10 | 11.550000 | 13.290000 | 0.49 | 82.260002 | 19.809999 |
| 999 | 1.98 | 0.45 | 2.12 | 6.270000 | 0.11 | 49.419998 | 75.730003 | 2.86 | 76.000000 | 25.160000 |

160 rows × 15 columns

In [102]:

```
x=b.iloc[:,0:10]
y=b.iloc[:, -1]
```

In [103]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [104]:

```
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[104]:

LinearRegression()

In [105]:

```
print(lr.intercept_)
```

28079025.314919464

In [106]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[106]:


| | Co-efficient |
|-------------|--------------|
| BEN | 10.322864 |
| CO | -96.321162 |
| EBE | -44.748442 |
| MXY | 26.519033 |
| NMHC | 346.911033 |
| NO_2 | 0.283597 |
| NOx | -0.083552 |
| OXY | -30.689805 |
| O_3 | 0.271343 |
| PM10 | 0.570190 |

In [107]:

```
prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[107]:

<matplotlib.collections.PathCollection at 0x1cb80829b20>



In [108]:

```
print(lr.score(x_test,y_test))
```

0.26497089777435334

In [109]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[109]:

Ridge(alpha=10)

In [110]:

```
rr.score(x_test,y_test)
```

Out[110]:

0.09013363420469

In [111]:

```
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[111]:

Lasso(alpha=10)

In [112]:

```
la.score(x_test,y_test)
```

Out[112]:

-0.016235272401284417

In [113]:

```
from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[113]:

ElasticNet()

In [114]:

```
print(en.coef_)
```

```
[-0.          -0.50619145 -3.18185688  0.           0.           1.06276317
 -0.28769662 -2.82256335  0.71368423  0.53976801]
```

In [115]:

```
print(en.intercept_)
```

28078969.129685692

In [116]:

```
prediction=en.predict(x_test)
print(en.score(x_test,y_test))
```

0.05127349309556084

In [117]:

```
f=StandardScaler().fit_transform(x)
```

In [118]:

```
logr=LogisticRegression()
logr.fit(f,y)
```

Out[118]:

LogisticRegression()

In [120]:

```
g=[[10,20,30,40,50,60,70,80,90,10]]
```

In [121]:

```
prediction=logr.predict(g)
print(prediction)
```

[28079006]

In [122]:

```
logr.classes_
```

Out[122]:

array([28079006, 28079024, 28079035, 28079099], dtype=int64)

In [123]:

```
logr.predict_proba(g)[0][0]
```

Out[123]:

1.0

In [131]:

```
logr.score(x_test,y_test)
```

Out[131]:

0.25

In [124]:

```
rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[124]:

```
RandomForestClassifier()
```

In [125]:

```
parameters={'max_depth':[1,2,3,4,5],  
            'min_samples_leaf':[5,10,15,20,25],  
            'n_estimators':[10,20,30,40,50]  
            }
```

In [126]:

```
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

Out[126]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
             param_grid={'max_depth': [1, 2, 3, 4, 5],  
                         'min_samples_leaf': [5, 10, 15, 20, 25],  
                         'n_estimators': [10, 20, 30, 40, 50]},  
             scoring='accuracy')
```

In [127]:

```
grid_search.best_score_
```

Out[127]:

```
0.6875
```

In [128]:

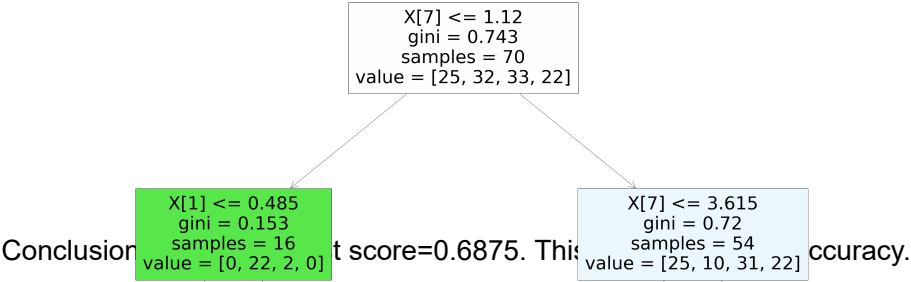
```
rfc_best=grid_search.best_estimator_
```

In [130]:

```
plt.figure(figsize=(80,80))
plot_tree(rfc_best.estimators_[5],filled=True)
```

Out[130]:

```
[Text(1785.6, 3986.4, 'X[7] <= 1.12\ngini = 0.743\nsamples = 70\nvalue =
[25, 32, 33, 22]'),
Text(892.8, 3261.6000000000004, 'X[1] <= 0.485\ngini = 0.153\nsamples = 1
6\nvalue = [0, 22, 2, 0]'),
Text(446.4, 2536.8, 'gini = 0.408\nsamples = 5\nvalue = [0, 5, 2, 0]'),
Text(1339.1999999999998, 2536.8, 'gini = 0.0\nsamples = 11\nvalue = [0, 1
7, 0, 0]'),
Text(2678.3999999999996, 3261.6000000000004, 'X[7] <= 3.615\ngini = 0.72
\nsamples = 54\nvalue = [25, 10, 31, 22]'),
Text(2232.0, 2536.8, 'X[4] <= 0.02\ngini = 0.699\nsamples = 45\nvalue =
[11, 10, 30, 22]'),
Text(1785.6, 1812.0, 'gini = 0.0\nsamples = 5\nvalue = [7, 0, 0, 0]'),
Text(2678.3999999999996, 1812.0, 'X[7] <= 1.555\ngini = 0.656\nsamples =
40\nvalue = [4, 10, 30, 22]'),
Text(1785.6, 1087.1999999999998, 'X[1] <= 0.725\ngini = 0.676\nsamples =
11\nvalue = [4, 7, 2, 2]'),
Text(1339.1999999999998, 362.39999999999964, 'gini = 0.625\nsamples = 6\n
value = [0, 4, 2, 2]'),
Text(2232.0, 362.39999999999964, 'gini = 0.49\nsamples = 5\nvalue = [4,
3, 0, 0]'),
Text(3571.2, 1087.1999999999998, 'X[5] <= 94.98\ngini = 0.541\nsamples =
29\nvalue = [0, 3, 28, 20]'),
Text(3124.7999999999997, 362.39999999999964, 'gini = 0.563\nsamples = 23
\nvalue = [0, 3, 16, 20]'),
Text(4017.6, 362.39999999999964, 'gini = 0.0\nsamples = 6\nvalue = [0, 0,
12, 0]'),
Text(3124.7999999999997, 2536.8, 'gini = 0.124\nsamples = 9\nvalue = [14,
0, 1, 0]')]
```



Madrid 2003

In [132]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\madrid_2003.csv")
a
```

Out[132]:

| | date | BEN | CO | EBE | MXV | NMHC | NO_2 | NOx | OXY | O_3 |
|--------|---------------------|------|------|------|------|------|-----------|------------|------|-----------|
| 0 | 2003-03-01 01:00:00 | NaN | 1.72 | NaN | NaN | NaN | 73.900002 | 316.299998 | NaN | 10.550000 |
| 1 | 2003-03-01 01:00:00 | NaN | 1.45 | NaN | NaN | 0.26 | 72.110001 | 250.000000 | 0.73 | 6.720000 |
| 2 | 2003-03-01 01:00:00 | NaN | 1.57 | NaN | NaN | NaN | 80.559998 | 224.199997 | NaN | 21.049999 |
| 3 | 2003-03-01 01:00:00 | NaN | 2.45 | NaN | NaN | NaN | 78.370003 | 450.399994 | NaN | 4.220000 |
| 4 | 2003-03-01 01:00:00 | NaN | 3.26 | NaN | NaN | NaN | 96.250000 | 479.100006 | NaN | 8.460000 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 243979 | 2003-10-01 00:00:00 | 0.20 | 0.16 | 2.01 | 3.17 | 0.02 | 31.799999 | 32.299999 | 1.68 | 34.049999 |
| 243980 | 2003-10-01 00:00:00 | 0.32 | 0.08 | 0.36 | 0.72 | NaN | 10.450000 | 14.760000 | 1.00 | 34.610001 |
| 243981 | 2003-10-01 00:00:00 | NaN | NaN | NaN | NaN | 0.07 | 34.639999 | 50.810001 | NaN | 32.160000 |
| 243982 | 2003-10-01 00:00:00 | NaN | NaN | NaN | NaN | 0.07 | 32.580002 | 41.020000 | NaN | NaN |
| 243983 | 2003-10-01 00:00:00 | 1.00 | 0.29 | 2.15 | 6.41 | 0.07 | 37.150002 | 56.849998 | 2.28 | 21.480000 |

243984 rows × 16 columns

In [133]:

```
a=a.head(2000)
a
```

Out[133]:

| | date | BEN | CO | EBE | MXV | NMHC | NO_2 | NOx | OXY | O_3 | P |
|------|---------------------|-----|------|-----|-----|------|-----------|------------|------|-----------|--------|
| 0 | 2003-03-01 01:00:00 | NaN | 1.72 | NaN | NaN | NaN | 73.900002 | 316.299988 | NaN | 10.550000 | 55.209 |
| 1 | 2003-03-01 01:00:00 | NaN | 1.45 | NaN | NaN | 0.26 | 72.110001 | 250.000000 | 0.73 | 6.720000 | 52.389 |
| 2 | 2003-03-01 01:00:00 | NaN | 1.57 | NaN | NaN | NaN | 80.559998 | 224.199997 | NaN | 21.049999 | 63.240 |
| 3 | 2003-03-01 01:00:00 | NaN | 2.45 | NaN | NaN | NaN | 78.370003 | 450.399994 | NaN | 4.220000 | 67.839 |
| 4 | 2003-03-01 01:00:00 | NaN | 3.26 | NaN | NaN | NaN | 96.250000 | 479.100006 | NaN | 8.460000 | 95.779 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1995 | 2003-03-04 00:00:00 | NaN | 1.70 | NaN | NaN | 0.27 | 76.070000 | 291.399994 | NaN | 6.970000 | 19.670 |
| 1996 | 2003-03-04 00:00:00 | NaN | 1.49 | NaN | NaN | NaN | 64.769997 | 204.600006 | NaN | 2.380000 | 32.419 |
| 1997 | 2003-03-04 00:00:00 | NaN | 1.71 | NaN | NaN | NaN | 51.099998 | 130.600006 | NaN | 7.810000 | 22.070 |
| 1998 | 2003-03-04 00:00:00 | NaN | 1.54 | NaN | NaN | 0.36 | 73.330002 | 287.600006 | NaN | 4.030000 | 39.180 |
| 1999 | 2003-03-04 00:00:00 | NaN | 0.86 | NaN | NaN | NaN | 58.360001 | 109.400002 | NaN | 8.860000 | 19.700 |

2000 rows × 16 columns



In [135]:

```
b=a.dropna()  
b
```

Out[135]:

| | date | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 | PM10 |
|------|---------------------|------|------|-------|-------|------|-----------|------------|-------|-------|-----------|
| 5 | 2003-03-01 01:00:00 | 8.41 | 1.94 | 9.83 | 21.49 | 0.45 | 90.300003 | 384.899994 | 9.48 | 9.95 | 95.150001 |
| 23 | 2003-03-01 01:00:00 | 3.46 | 1.27 | 3.43 | 7.08 | 0.18 | 54.250000 | 173.300003 | 3.37 | 6.54 | 53.009999 |
| 27 | 2003-03-01 01:00:00 | 6.39 | 1.79 | 5.75 | 10.88 | 0.33 | 75.459999 | 281.100006 | 3.68 | 6.69 | 63.840001 |
| 33 | 2003-03-01 02:00:00 | 7.42 | 1.47 | 10.63 | 24.73 | 0.35 | 83.309998 | 277.200012 | 11.00 | 9.90 | 58.880001 |
| 51 | 2003-03-01 02:00:00 | 3.62 | 1.29 | 3.20 | 7.08 | 0.19 | 42.209999 | 166.300003 | 3.41 | 6.38 | 47.599999 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1959 | 2003-03-03 22:00:00 | 3.38 | 1.25 | 3.17 | 6.34 | 0.23 | 60.830002 | 172.500000 | 2.66 | 7.19 | 29.990001 |
| 1965 | 2003-03-03 23:00:00 | 5.59 | 1.19 | 7.04 | 17.40 | 0.28 | 62.180000 | 200.800003 | 8.45 | 10.09 | 19.299999 |
| 1983 | 2003-03-03 23:00:00 | 1.10 | 0.64 | 0.84 | 1.75 | 0.07 | 51.970001 | 83.900002 | 1.05 | 5.91 | 35.009999 |
| 1987 | 2003-03-03 23:00:00 | 3.76 | 1.20 | 3.54 | 7.02 | 0.23 | 60.580002 | 170.100006 | 2.86 | 6.71 | 26.620001 |
| 1993 | 2003-03-04 00:00:00 | 4.54 | 1.69 | 6.07 | 15.06 | 0.35 | 71.220001 | 314.700012 | 7.48 | 10.07 | 46.349999 |

212 rows × 16 columns

In [136]:

```
b.columns
```

Out[136]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
      'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],  
      dtype='object')
```

In [137]:

```
b=b[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
      'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
b
```

Out[137]:

| | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 | PM10 | PXY |
|-------------|------|------|-------|-------|------|-----------|------------|-------|-------|-----------|------|
| 5 | 8.41 | 1.94 | 9.83 | 21.49 | 0.45 | 90.300003 | 384.899994 | 9.48 | 9.95 | 95.150002 | 7.94 |
| 23 | 3.46 | 1.27 | 3.43 | 7.08 | 0.18 | 54.250000 | 173.300003 | 3.37 | 6.54 | 53.009998 | 2.62 |
| 27 | 6.39 | 1.79 | 5.75 | 10.88 | 0.33 | 75.459999 | 281.100006 | 3.68 | 6.69 | 63.840000 | 4.24 |
| 33 | 7.42 | 1.47 | 10.63 | 24.73 | 0.35 | 83.309998 | 277.200012 | 11.00 | 9.90 | 58.880001 | 8.93 |
| 51 | 3.62 | 1.29 | 3.20 | 7.08 | 0.19 | 42.209999 | 166.300003 | 3.41 | 6.38 | 47.599998 | 2.70 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1959 | 3.38 | 1.25 | 3.17 | 6.34 | 0.23 | 60.830002 | 172.500000 | 2.66 | 7.19 | 29.990000 | 2.99 |
| 1965 | 5.59 | 1.19 | 7.04 | 17.40 | 0.28 | 62.180000 | 200.800003 | 8.45 | 10.09 | 19.299999 | 6.90 |
| 1983 | 1.10 | 0.64 | 0.84 | 1.75 | 0.07 | 51.970001 | 83.900002 | 1.05 | 5.91 | 35.009998 | 0.72 |
| 1987 | 3.76 | 1.20 | 3.54 | 7.02 | 0.23 | 60.580002 | 170.100006 | 2.86 | 6.71 | 26.620001 | 3.28 |
| 1993 | 4.54 | 1.69 | 6.07 | 15.06 | 0.35 | 71.220001 | 314.700012 | 7.48 | 10.07 | 46.349998 | 5.99 |

212 rows × 15 columns

In [138]:

```
x=b.iloc[:,0:5]
y=b.iloc[:, -1]
```

In [139]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [140]:

```
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[140]:

LinearRegression()

In [141]:

```
print(lr.intercept_)
```

28079100.77042889

In [142]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[142]:


| | Co-efficient |
|-------------|--------------|
| BEN | -4.354966 |
| CO | -170.372248 |
| EBE | 16.835167 |
| MXY | -11.874334 |
| NMHC | 860.450085 |

In [143]:

```
prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[143]:

<matplotlib.collections.PathCollection at 0x1cb80a0d790>



In [144]:

```
print(lr.score(x_test,y_test))
```

0.6422489114361061

In [145]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[145]:

Ridge(alpha=10)

In [146]:

```
rr.score(x_test,y_test)
```

Out[146]:

0.1877820357689176

In [147]:

```
la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

Out[147]:

Lasso(alpha=10)

In [148]:

```
la.score(x_test,y_test)
```

Out[148]:

0.04866758105027491

In [149]:

```
en=ElasticNet()  
en.fit(x_train,y_train)
```

Out[149]:

ElasticNet()

In [150]:

```
print(en.coef_)
```

[-1.06400298 1.31833139 3.55029825 -4.02609241 1.16258014]

In [151]:

```
print(en.intercept_)
```

28079055.404426787

In [152]:

```
prediction=en.predict(x_test)  
print(en.score(x_test,y_test))
```

0.07228820797072222

In [153]:

```
f=StandardScaler().fit_transform(x)
```

In [154]:

```
logr=LogisticRegression()  
logr.fit(f,y)
```

Out[154]:

```
LogisticRegression()
```

In [155]:

```
g=[[10,20,30,40,50]]
```

In [156]:

```
prediction=logr.predict(g)  
print(prediction)
```

```
[28079099]
```

In [157]:

```
logr.classes_
```

Out[157]:

```
array([28079006, 28079024, 28079099], dtype=int64)
```

In [158]:

```
logr.predict_proba(g)[0][0]
```

Out[158]:

```
1.3162606173431188e-26
```

In [159]:

```
logr.score(x_test,y_test)
```

Out[159]:

```
0.34375
```

In [160]:

```
rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[160]:

```
RandomForestClassifier()
```

In [161]:

```
parameters={'max_depth':[1,2,3,4,5],  
            'min_samples_leaf':[5,10,15,20,25],  
            'n_estimators':[10,20,30,40,50]  
            }
```

In [162]:

```
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

Out[162]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),
             param_grid={'max_depth': [1, 2, 3, 4, 5],
                          'min_samples_leaf': [5, 10, 15, 20, 25],
                          'n_estimators': [10, 20, 30, 40, 50]},
             scoring='accuracy')
```

In [163]:

```
grid_search.best_score_
```

Out[163]:

```
0.8040540540540541
```

In [164]:

```
rfc_best=grid_search.best_estimator_
```

In [165]:

```
plt.figure(figsize=(80,80))
plot_tree(rfc_best.estimators_[5],filled=True)
```

Out[165]:

```
[Text(2575.3846153846152, 3805.2000000000003, 'X[3] <= 4.91\ngini = 0.662\nsamples = 91\nvalue = [55, 52, 41]'),
Text(1373.5384615384614, 2718.0, 'X[0] <= 1.765\ngini = 0.553\nsamples = 56\nvalue = [7, 50, 34]'),
Text(686.7692307692307, 1630.8000000000002, 'X[0] <= 1.16\ngini = 0.461\nsamples = 27\nvalue = [3, 10, 29]'),
Text(343.38461538461536, 543.5999999999999, 'gini = 0.569\nsamples = 10\nvalue = [3, 7, 2]'),
Text(1030.1538461538462, 543.5999999999999, 'gini = 0.18\nsamples = 17\nvalue = [0, 3, 27]'),
Text(2060.3076923076924, 1630.8000000000002, 'X[1] <= 0.545\ngini = 0.317\nsamples = 29\nvalue = [4, 40, 5]'),
Text(1716.9230769230767, 543.5999999999999, 'gini = 0.0\nsamples = 18\nvalue = [0, 32, 0]'),
Text(2403.6923076923076, 543.5999999999999, 'gini = 0.637\nsamples = 11\nvalue = [4, 8, 5]'),
Text(3777.230769230769, 2718.0, 'X[3] <= 8.375\ngini = 0.275\nsamples = 35\nvalue = [48, 2, 7]'),
Text(3433.8461538461534, 1630.8000000000002, 'X[2] <= 2.62\ngini = 0.461\nsamples = 20\nvalue = [20, 2, 7]'),
Text(3090.461538461538, 543.5999999999999, 'gini = 0.105\nsamples = 10\nvalue = [17, 1, 0]'),
Text(3777.230769230769, 543.5999999999999, 'gini = 0.512\nsamples = 10\nvalue = [3, 1, 7]'),
Text(4120.615384615385, 1630.8000000000002, 'gini = 0.0\nsamples = 15\nvalue = [28, 0, 0]')]
```

Conclusion: RandomForest Score=0.8040540540540541. It has the highest accuracy.

Madrid 2004

In [166]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\madrid_2004.csv")
a
```

Out[166]:

| | | date | BEN | CO | EBE | MXV | NMHC | NO_2 | NOx | OXY | O_3 |
|--------|---------------------|------|------|------|------|------|------------|------------|------------|-----------|-----------|
| | | | | | | | | | | | |
| 0 | 2004-08-01 01:00:00 | NaN | 0.66 | NaN | NaN | NaN | NaN | 89.550003 | 118.900002 | NaN | 40.020000 |
| 1 | 2004-08-01 01:00:00 | 2.66 | 0.54 | 2.99 | 6.08 | 0.18 | 51.799999 | 53.860001 | 3.28 | 51.689999 | 22 |
| 3 | 2004-08-01 01:00:00 | NaN | 0.53 | NaN | NaN | NaN | NaN | 87.290001 | 105.000000 | NaN | 36.730000 |
| 4 | 2004-08-01 01:00:00 | NaN | 0.17 | NaN | NaN | NaN | NaN | 34.910000 | 35.349998 | NaN | 86.269997 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 245491 | 2004-06-01 00:00:00 | 0.75 | 0.21 | 0.85 | 1.55 | 0.07 | 59.580002 | 64.389999 | 0.66 | 33.029999 | 30 |
| 245492 | 2004-06-01 00:00:00 | 2.49 | 0.75 | 2.44 | 4.57 | NaN | 97.139999 | 146.899994 | 2.34 | 7.740000 | 37 |
| 245493 | 2004-06-01 00:00:00 | NaN | NaN | NaN | NaN | 0.13 | 102.699997 | 132.600006 | NaN | 17.809999 | 22 |
| 245494 | 2004-06-01 00:00:00 | NaN | NaN | NaN | NaN | 0.09 | 82.599998 | 102.599998 | NaN | NaN | 45 |
| 245495 | 2004-06-01 00:00:00 | 3.01 | 0.67 | 2.78 | 5.12 | 0.20 | 92.550003 | 141.000000 | 2.60 | 11.460000 | 24 |

245496 rows × 17 columns

In [167]:

```
a=a.head(2000)
a
```

Out[167]:

| | date | BEN | CO | EBE | MXV | NMHC | NO_2 | NOx | OXY | O_3 | P |
|------|---------------------|------|------|------|------|------|-----------|------------|------|-----------|--------|
| 0 | 2004-08-01 01:00:00 | NaN | 0.66 | NaN | NaN | NaN | 89.550003 | 118.900002 | NaN | 40.020000 | 39.990 |
| 1 | 2004-08-01 01:00:00 | 2.66 | 0.54 | 2.99 | 6.08 | 0.18 | 51.799999 | 53.860001 | 3.28 | 51.689999 | 22.950 |
| 2 | 2004-08-01 01:00:00 | NaN | 1.02 | NaN | NaN | NaN | 93.389999 | 138.600006 | NaN | 20.860001 | 49.480 |
| 3 | 2004-08-01 01:00:00 | NaN | 0.53 | NaN | NaN | NaN | 87.290001 | 105.000000 | NaN | 36.730000 | 31.070 |
| 4 | 2004-08-01 01:00:00 | NaN | 0.17 | NaN | NaN | NaN | 34.910000 | 35.349998 | NaN | 86.269997 | 54.080 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1995 | 2004-08-04 01:00:00 | NaN | 0.24 | NaN | NaN | NaN | 33.070000 | 36.939999 | NaN | 38.599998 | 4.190 |
| 1996 | 2004-08-04 01:00:00 | NaN | 0.02 | NaN | NaN | NaN | 18.049999 | 19.959999 | NaN | 60.380001 | 1.940 |
| 1997 | 2004-08-04 01:00:00 | 0.93 | 0.24 | 0.94 | 1.50 | 0.01 | 29.910000 | 40.490002 | 1.25 | 50.660000 | 3.900 |
| 1998 | 2004-08-04 01:00:00 | NaN | 0.26 | NaN | NaN | 0.09 | 29.660000 | 35.599998 | NaN | 49.730000 | 5.320 |
| 1999 | 2004-08-04 01:00:00 | 0.23 | 0.33 | 0.90 | NaN | 0.00 | 47.259998 | 60.220001 | NaN | 33.980000 | 9.520 |

2000 rows × 17 columns



In [168]:

```
b=a.dropna()  
b
```

Out[168]:

| | date | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 | |
|------|---------------------|------|------|------|------|------|------------|------------|------|-----------|------|
| 5 | 2004-08-01 01:00:00 | 3.24 | 0.63 | 5.55 | 9.72 | 0.06 | 103.800003 | 144.800003 | 5.04 | 32.480000 | 59.1 |
| 22 | 2004-08-01 01:00:00 | 0.55 | 0.36 | 0.54 | 0.86 | 0.07 | 31.980000 | 32.799999 | 0.50 | 79.040001 | 43.5 |
| 26 | 2004-08-01 01:00:00 | 1.80 | 0.46 | 2.28 | 4.62 | 0.21 | 62.259998 | 75.470001 | 2.47 | 54.419998 | 46.6 |
| 32 | 2004-08-01 02:00:00 | 1.94 | 0.67 | 3.14 | 4.91 | 0.06 | 113.500000 | 165.800003 | 2.56 | 26.980000 | 86.9 |
| 49 | 2004-08-01 02:00:00 | 0.29 | 0.30 | 0.47 | 0.76 | 0.07 | 33.919998 | 34.840000 | 0.46 | 75.570000 | 48.9 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1963 | 2004-08-03 23:00:00 | 1.23 | 0.39 | 1.34 | 3.08 | 0.11 | 52.779999 | 68.750000 | 2.79 | 32.990002 | 30.2 |
| 1969 | 2004-08-04 00:00:00 | 1.02 | 0.36 | 1.00 | 1.18 | 0.02 | 43.529999 | 66.279999 | 1.55 | 42.520000 | 25.5 |
| 1987 | 2004-08-04 00:00:00 | 0.58 | 0.08 | 0.51 | 0.64 | 0.00 | 10.330000 | 10.820000 | 0.68 | 46.009998 | 16.4 |
| 1991 | 2004-08-04 00:00:00 | 1.32 | 0.27 | 1.31 | 1.90 | 0.08 | 35.349998 | 46.880001 | 1.43 | 39.380001 | 32.4 |
| 1997 | 2004-08-04 01:00:00 | 0.93 | 0.24 | 0.94 | 1.50 | 0.01 | 29.910000 | 40.490002 | 1.25 | 50.660000 | 3.9 |

195 rows × 17 columns

In [169]:

```
b.columns
```

Out[169]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
      'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],  
      dtype='object')
```

In [170]:

```
b=b[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
      'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
b
```

Out[170]:

| | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 | PM10 | PM25 | PXY |
|------|------|------|------|------|------|------------|------------|------|-----------|-----------|-----------|------|
| 5 | 3.24 | 0.63 | 5.55 | 9.72 | 0.06 | 103.800003 | 144.800003 | 5.04 | 32.480000 | 59.110001 | 38.049999 | 4.16 |
| 22 | 0.55 | 0.36 | 0.54 | 0.86 | 0.07 | 31.980000 | 32.799999 | 0.50 | 79.040001 | 43.549999 | 22.780001 | 0.39 |
| 26 | 1.80 | 0.46 | 2.28 | 4.62 | 0.21 | 62.259998 | 75.470001 | 2.47 | 54.419998 | 46.630001 | 29.459999 | 2.21 |
| 32 | 1.94 | 0.67 | 3.14 | 4.91 | 0.06 | 113.500000 | 165.800003 | 2.56 | 26.980000 | 86.930000 | 45.639999 | 2.14 |
| 49 | 0.29 | 0.30 | 0.47 | 0.76 | 0.07 | 33.919998 | 34.840000 | 0.46 | 75.570000 | 48.959999 | 25.959999 | 0.34 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1963 | 1.23 | 0.39 | 1.34 | 3.08 | 0.11 | 52.779999 | 68.750000 | 2.79 | 32.990002 | 30.200001 | 19.670000 | 2.73 |
| 1969 | 1.02 | 0.36 | 1.00 | 1.18 | 0.02 | 43.529999 | 66.279999 | 1.55 | 42.520000 | 25.570000 | 12.660000 | 1.72 |
| 1987 | 0.58 | 0.08 | 0.51 | 0.64 | 0.00 | 10.330000 | 10.820000 | 0.68 | 46.009998 | 16.410000 | 6.300000 | 0.27 |
| 1994 | 1.23 | 0.37 | 1.31 | 1.00 | 0.08 | 35.340000 | 46.880001 | 1.42 | 30.380001 | 22.420000 | 12.360000 | 1.28 |

In [171]:

```
x=b.iloc[:,0:6]
y=b.iloc[:, -1]
```

In [172]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [173]:

```
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[173]:

LinearRegression()

In [174]:

```
print(lr.score(x_test,y_test))
```

0.18735247964299628

In [175]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[175]:

Ridge(alpha=10)

In [176]:

```
rr.score(x_test,y_test)
```

Out[176]:

0.11385632125962974

In [177]:

```
la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

Out[177]:

Lasso(alpha=10)

In [178]:

```
la.score(x_test,y_test)
```

Out[178]:

0.018612493009876552

In [179]:

```
en=ElasticNet()  
en.fit(x_train,y_train)
```

Out[179]:

ElasticNet()

In [180]:

```
prediction=en.predict(x_test)  
print(en.score(x_test,y_test))
```

0.034590122360959485

In [181]:

```
f=StandardScaler().fit_transform(x)
```

In [182]:

```
logr=LogisticRegression()  
logr.fit(f,y)
```

Out[182]:

LogisticRegression()

In [183]:

```
g=[[10,20,30,40,50,60]]
```

In [184]:

```
prediction=logr.predict(g)
print(prediction)
```

[28079006]

In [185]:

```
logr.predict_proba(g)[0][0]
```

Out[185]:

1.0

In [186]:

```
logr.score(x_test,y_test)
```

Out[186]:

0.2542372881355932

In [187]:

```
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[187]:

RandomForestClassifier()

In [188]:

```
parameters={'max_depth':[1,2,3,4,5],
            'min_samples_leaf':[5,10,15,20,25],
            'n_estimators':[10,20,30,40,50]}
}
```

In [189]:

```
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

Out[189]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),
             param_grid={'max_depth': [1, 2, 3, 4, 5],
                         'min_samples_leaf': [5, 10, 15, 20, 25],
                         'n_estimators': [10, 20, 30, 40, 50]},
             scoring='accuracy')
```

In [190]:

```
grid_search.best_score_
```

Out[190]:

0.9044117647058824

In [191]:

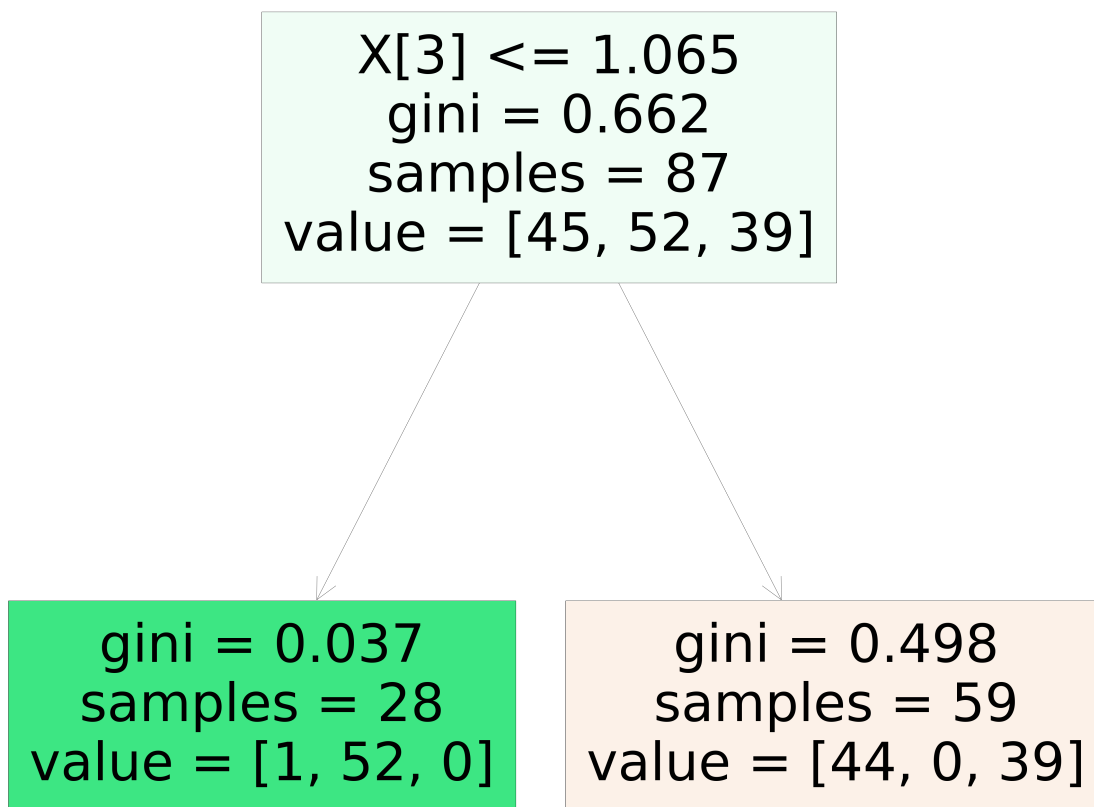
```
rfc_best=grid_search.best_estimator_
```

In [192]:

```
plt.figure(figsize=(80,80))  
plot_tree(rfc_best.estimators_[5],filled=True)
```

Out[192]:

```
[Text(2232.0, 3261.6000000000004, 'X[3] <= 1.065\nngini = 0.662\nsamples =  
87\nvalue = [45, 52, 39]'),  
Text(1116.0, 1087.1999999999998, 'gini = 0.037\nsamples = 28\nvalue = [1,  
52, 0]'),  
Text(3348.0, 1087.1999999999998, 'gini = 0.498\nsamples = 59\nvalue = [4  
4, 0, 39]')]
```



Conclusion: RandomForest Score=0.9044117647058824. It has the highest accuracy.