

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
import re
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
```

Madrid 2009

In [2]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\madrid_2009.csv")
a
```

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	
0	2009-10-01 01:00:00	NaN	0.27	NaN	NaN	NaN	39.889999	48.150002	NaN	50.680000	18.0
1	2009-10-01 01:00:00	NaN	0.22	NaN	NaN	NaN	21.230000	24.260000	NaN	55.880001	10.0
2	2009-10-01 01:00:00	NaN	0.18	NaN	NaN	NaN	31.230000	34.880001	NaN	49.060001	25.0
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	26.0
4	2009-10-01 01:00:00	NaN	0.41	NaN	NaN	0.12	61.349998	76.260002	NaN	38.090000	23.0
...
215683	2009-06-01 00:00:00	0.50	0.22	0.39	0.75	0.09	22.000000	24.510000	1.00	82.239998	10.0
215684	2009-06-01 00:00:00	NaN	0.31	NaN	NaN	NaN	76.110001	101.099998	NaN	41.220001	9.0
215685	2009-06-01 00:00:00	0.13	NaN	0.86	NaN	0.23	81.050003	99.849998	NaN	24.830000	12.0
215686	2009-06-01 00:00:00	0.21	NaN	2.96	NaN	0.10	72.419998	82.959999	NaN	NaN	13.0
215687	2009-06-01 00:00:00	0.37	0.32	0.99	1.36	0.14	54.290001	64.480003	1.06	56.919998	15.0

215688 rows × 17 columns



In [3]:

```
a.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 215688 entries, 0 to 215687
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        215688 non-null object  
 1   BEN         60082 non-null  float64
 2   CO          190801 non-null float64
 3   EBE         60081 non-null  float64
 4   MXY         24846 non-null  float64
 5   NMHC        74748 non-null  float64
 6   NO_2        214562 non-null float64
 7   NOx         214565 non-null float64
 8   OXY         24854 non-null  float64
 9   O_3         204482 non-null float64
10  PM10        196331 non-null float64
11  PM25        55822 non-null  float64
12  PXY         24854 non-null  float64
13  SO_2        212671 non-null float64
14  TCH         75213 non-null  float64
15  TOL         59920 non-null  float64
16  station     215688 non-null int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 28.0+ MB
```

In [4]:

```
b=a.fillna(value=87)
b
```

Out[4]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
0	2009-10-01 01:00:00	87.00	0.27	87.00	87.00	87.00	39.889999	48.150002	87.00	50.680000
1	2009-10-01 01:00:00	87.00	0.22	87.00	87.00	87.00	21.230000	24.260000	87.00	55.880001
2	2009-10-01 01:00:00	87.00	0.18	87.00	87.00	87.00	31.230000	34.880001	87.00	49.060001
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998
4	2009-10-01 01:00:00	87.00	0.41	87.00	87.00	0.12	61.349998	76.260002	87.00	38.090000
...
215683	2009-06-01 00:00:00	0.50	0.22	0.39	0.75	0.09	22.000000	24.510000	1.00	82.239998
215684	2009-06-01 00:00:00	87.00	0.31	87.00	87.00	87.00	76.110001	101.099998	87.00	41.220001
215685	2009-06-01 00:00:00	0.13	87.00	0.86	87.00	0.23	81.050003	99.849998	87.00	24.830000
215686	2009-06-01 00:00:00	0.21	87.00	2.96	87.00	0.10	72.419998	82.959999	87.00	87.000000
215687	2009-06-01 00:00:00	0.37	0.32	0.99	1.36	0.14	54.290001	64.480003	1.06	56.919998

215688 rows × 17 columns

In [5]:

```
b.columns
```

Out[5]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
      'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

In [6]:

```
c=b.head(10)
c
```

Out[6]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM
0	2009-10-01 01:00:00	87.00	0.27	87.00	87.00	87.00	39.889999	48.150002	87.00	50.680000	18.260000
1	2009-10-01 01:00:00	87.00	0.22	87.00	87.00	87.00	21.230000	24.260000	87.00	55.880001	10.580000
2	2009-10-01 01:00:00	87.00	0.18	87.00	87.00	87.00	31.230000	34.880001	87.00	49.060001	25.190000
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	26.530000
4	2009-10-01 01:00:00	87.00	0.41	87.00	87.00	0.12	61.349998	76.260002	87.00	38.090000	23.760000
5	2009-10-01 01:00:00	87.00	0.29	87.00	87.00	87.00	43.200001	50.080002	87.00	35.840000	21.870000
6	2009-10-01 01:00:00	87.00	0.20	87.00	87.00	87.00	35.430000	38.520000	87.00	33.549999	17.350000
7	2009-10-01 01:00:00	87.00	0.15	87.00	87.00	87.00	27.309999	33.150002	87.00	53.549999	16.520000
8	2009-10-01 01:00:00	87.00	0.21	87.00	87.00	0.39	33.889999	40.799999	87.00	58.549999	16.650000
9	2009-10-01 01:00:00	87.00	0.32	87.00	87.00	87.00	46.349998	60.540001	87.00	45.340000	15.160000

In [7]:

```
d=c[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
      'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
d
```

Out[7]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PXY
0	87.00	0.27	87.00	87.00	87.00	39.889999	48.150002	87.00	50.680000	18.260000	87.0
1	87.00	0.22	87.00	87.00	87.00	21.230000	24.260000	87.00	55.880001	10.580000	87.0
2	87.00	0.18	87.00	87.00	87.00	31.230000	34.880001	87.00	49.060001	25.190001	87.0
3	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	26.530001	1.3
4	87.00	0.41	87.00	87.00	0.12	61.349998	76.260002	87.00	38.090000	23.760000	87.0
5	87.00	0.29	87.00	87.00	87.00	43.200001	50.080002	87.00	35.840000	21.870001	87.0
6	87.00	0.20	87.00	87.00	87.00	35.430000	38.520000	87.00	33.549999	17.350000	87.0
7	87.00	0.15	87.00	87.00	87.00	27.309999	33.150002	87.00	53.549999	16.520000	87.0
8	87.00	0.21	87.00	87.00	0.39	33.889999	40.799999	87.00	58.549999	16.650000	87.0
9	87.00	0.32	87.00	87.00	87.00	46.349998	60.540001	87.00	45.340000	15.160000	87.0

In [8]:

```
x=d[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY']]
y=d['TCH']
```

In [9]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [10]:

```
from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[10]:

LinearRegression()

In [11]:

```
print(lr.intercept_)
0.6181014730252841
```

In [12]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[12]:


	Co-efficient
BEN	1.398881e-13
CO	-3.428767e+00
EBE	-6.342149e-15
MXY	-4.440892e-16
NMHC	9.878962e-01
NO_2	2.564760e-02
NOx	6.348843e-03
OXY	0.000000e+00

In [13]:

```
prediction=lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[13]:

<matplotlib.collections.PathCollection at 0x2a95bf6d640>



In [14]:

```
print(lr.score(x_test,y_test))
```

0.9999338807553346

In [15]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [16]:

```
rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[16]:

Ridge(alpha=10)

In [17]:

```
rr.score(x_test,y_test)
```

Out[17]:

0.9999954967522613

In [18]:

```
la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

Out[18]:

Lasso(alpha=10)

In [19]:

```
la.score(x_test,y_test)
```

Out[19]:

0.9999534400930205

In [20]:

```
a1=b.head(7000)
a1
```

Out[20]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	
0	2009-10-01 01:00:00	87.00	0.27	87.00	87.00	87.00	39.889999	48.150002	87.00	50.680000	18
1	2009-10-01 01:00:00	87.00	0.22	87.00	87.00	87.00	21.230000	24.260000	87.00	55.880001	10
2	2009-10-01 01:00:00	87.00	0.18	87.00	87.00	87.00	31.230000	34.880001	87.00	49.060001	25
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	26
4	2009-10-01 01:00:00	87.00	0.41	87.00	87.00	0.12	61.349998	76.260002	87.00	38.090000	23
...	
6995	2009-10-12 16:00:00	0.42	0.74	0.43	1.08	0.49	11.680000	15.810000	0.67	84.389999	11
6996	2009-10-12 16:00:00	87.00	0.23	87.00	87.00	87.00	33.090000	54.380001	87.00	57.480000	16
6997	2009-10-12 16:00:00	0.13	87.00	0.31	87.00	0.19	27.670000	36.860001	87.00	56.240002	19
6998	2009-10-12 16:00:00	0.20	87.00	1.00	87.00	0.13	16.459999	30.200001	87.00	87.000000	30
6999	2009-10-12 16:00:00	0.23	0.25	0.63	1.08	0.18	22.760000	32.700001	0.67	64.739998	11

7000 rows × 17 columns



In [21]:

```
e=a1[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
      'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
e
```

Out[21]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10
0	87.00	0.27	87.00	87.00	87.00	39.889999	48.150002	87.00	50.680000	18.260000
1	87.00	0.22	87.00	87.00	87.00	21.230000	24.260000	87.00	55.880001	10.580000
2	87.00	0.18	87.00	87.00	87.00	31.230000	34.880001	87.00	49.060001	25.190001
3	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	26.530001
4	87.00	0.41	87.00	87.00	0.12	61.349998	76.260002	87.00	38.090000	23.760000
...
6995	0.42	0.74	0.43	1.08	0.49	11.680000	15.810000	0.67	84.389999	11.110000
6996	87.00	0.23	87.00	87.00	87.00	33.090000	54.380001	87.00	57.480000	16.969999
6997	0.13	87.00	0.31	87.00	0.19	27.670000	36.860001	87.00	56.240002	19.820000
6998	0.20	87.00	1.00	87.00	0.13	16.459999	30.200001	87.00	87.000000	30.650000
6999	0.23	0.25	0.63	1.08	0.18	22.760000	32.700001	0.67	64.739998	11.070000

7000 rows × 15 columns

In [22]:

```
f=e.iloc[:,0:14]
g=e.iloc[:, -1]
```

In [23]:

```
h=StandardScaler().fit_transform(f)
```

In [24]:

```
logr=LogisticRegression(max_iter=10000)
logr.fit(h,g)
```

Out[24]:

LogisticRegression(max_iter=10000)

In [25]:

```
from sklearn.model_selection import train_test_split
h_train,h_test,g_train,g_test=train_test_split(h,g,test_size=0.3)
```

In [26]:

```
i=[[10,20,30,40,50,60,11,22,33,44,55,54,21,78]]
```

In [27]:

```
prediction=logr.predict(i)  
print(prediction)
```

```
[28079021]
```

In [28]:

```
logr.classes_
```

Out[28]:

```
array([28079003, 28079004, 28079006, 28079007, 28079008, 28079009,  
       28079011, 28079012, 28079014, 28079016, 28079017, 28079018,  
       28079019, 28079021, 28079022, 28079023, 28079024, 28079025,  
       28079026, 28079027, 28079036, 28079038, 28079039, 28079040,  
       28079099], dtype=int64)
```

In [29]:

```
logr.predict_proba(i)[0][0]
```

Out[29]:

```
5.400803276887222e-153
```

In [30]:

```
logr.predict_proba(i)[0][1]
```

Out[30]:

```
5.714501500703068e-41
```

In [31]:

```
logr.score(h_test,g_test)
```

Out[31]:

```
0.5633333333333334
```

In [32]:

```
from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

Out[32]:

```
ElasticNet()
```

In [33]:

```
print(en.coef_)  
  
[0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00  
 9.86230309e-01 0.00000000e+00 5.33110402e-04 0.00000000e+00]
```

In [34]:

```
print(en.intercept_)  
  
1.1573688835558116
```

In [35]:

```
prediction=en.predict(x_test)  
print(en.score(x_test,y_test))  
  
0.9999988176366412
```

In [36]:

```
from sklearn.ensemble import RandomForestClassifier  
  
rfc=RandomForestClassifier()  
rfc.fit(h_train,g_train)
```

Out[36]:

```
RandomForestClassifier()
```

In [37]:

```
parameters={'max_depth':[1,2,3,4,5],  
            'min_samples_leaf':[5,10,15,20,25],  
            'n_estimators':[10,20,30,40,50]  
            }
```

In [38]:

```
from sklearn.model_selection import GridSearchCV  
  
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(h_train,g_train)
```

Out[38]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
             param_grid={'max_depth': [1, 2, 3, 4, 5],  
                         'min_samples_leaf': [5, 10, 15, 20, 25],  
                         'n_estimators': [10, 20, 30, 40, 50]},  
             scoring='accuracy')
```

In [39]:

```
grid_search.best_score_
```

Out[39]:

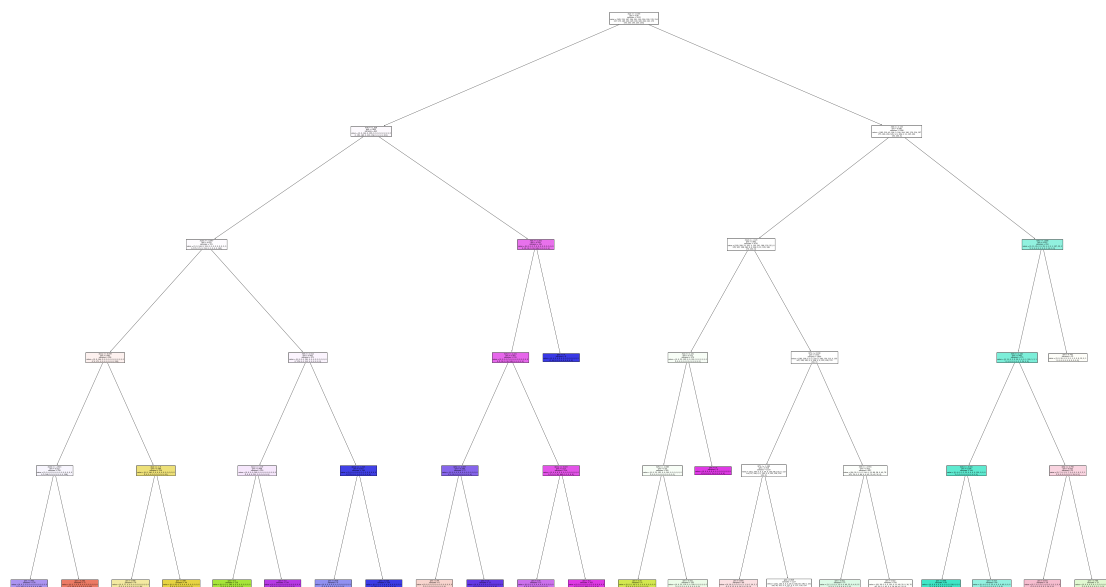
```
0.5775510204081633
```

In [40]:

```
rfc_best=grid_search.best_estimator_
```

In [41]:

```
from sklearn.tree import plot_tree  
  
plt.figure(figsize=(80,50))  
plot_tree(rfc_best.estimators_[2],filled=True)
```



Conclusion: Ridge score=0.9999960729058915.It has the highest accuracy.

Madrid 2010

In [42]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\madrid_2010.csv")
a
```

Out[42]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3
0	2010-03-01 01:00:00	NaN	0.29	NaN	NaN	NaN	25.090000	29.219999	NaN	68.930000
1	2010-03-01 01:00:00	NaN	0.27	NaN	NaN	NaN	24.879999	30.040001	NaN	NaN
2	2010-03-01 01:00:00	NaN	0.28	NaN	NaN	NaN	17.410000	20.540001	NaN	72.120003
3	2010-03-01 01:00:00	0.38	0.24	1.74	NaN	0.05	15.610000	21.080000	NaN	72.970001
4	2010-03-01 01:00:00	0.79	NaN	1.32	NaN	NaN	21.430000	26.070000	NaN	NaN
...
209443	2010-08-01 00:00:00	NaN	0.55	NaN	NaN	NaN	125.000000	219.899994	NaN	25.379999
209444	2010-08-01 00:00:00	NaN	0.27	NaN	NaN	NaN	45.709999	47.410000	NaN	NaN
209445	2010-08-01 00:00:00	NaN	NaN	NaN	NaN	0.24	46.560001	49.040001	NaN	46.250000
209446	2010-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	46.770000	50.119999	NaN	77.709999
209447	2010-08-01 00:00:00	0.92	0.43	0.71	NaN	0.25	76.330002	88.190002	NaN	52.259998

209448 rows × 17 columns



In [43]:

```
a=a.head(1000)
a
```

Out[43]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM
0	2010-03-01 01:00:00	NaN	0.29	NaN	NaN	NaN	25.090000	29.219999	NaN	68.930000	NaN
1	2010-03-01 01:00:00	NaN	0.27	NaN	NaN	NaN	24.879999	30.040001	NaN	NaN	NaN
2	2010-03-01 01:00:00	NaN	0.28	NaN	NaN	NaN	17.410000	20.540001	NaN	72.120003	NaN
3	2010-03-01 01:00:00	0.38	0.24	1.74	NaN	0.05	15.610000	21.080000	NaN	72.970001	19.410000
4	2010-03-01 01:00:00	0.79	NaN	1.32	NaN	NaN	21.430000	26.070000	NaN	NaN	24.670000
...
995	2010-03-02 18:00:00	0.51	0.20	0.91	1.27	0.39	20.330000	22.940001	1.42	86.410004	14.280000
996	2010-03-02 18:00:00	NaN	NaN	NaN	NaN	0.13	28.370001	40.669998	NaN	73.480003	NaN
997	2010-03-02 18:00:00	NaN	NaN	NaN	NaN	NaN	44.029999	50.509998	NaN	NaN	22.049999
998	2010-03-02 18:00:00	NaN	NaN	NaN	NaN	NaN	31.770000	37.040001	NaN	85.040001	NaN
999	2010-03-02 18:00:00	NaN	NaN	NaN	NaN	NaN	32.500000	39.279999	NaN	NaN	16.350000

1000 rows × 17 columns



In [44]:

```
b=a.dropna()  
b
```

Out[44]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10
11	2010-03-01 01:00:00	0.78	0.18	0.84	0.73	0.28	10.420000	11.900000	1.00	90.309998	18.370000
23	2010-03-01 01:00:00	0.70	0.23	1.00	0.73	0.18	17.820000	22.290001	1.00	70.550003	23.639999
35	2010-03-01 02:00:00	0.58	0.17	0.84	0.73	0.28	3.500000	4.950000	1.00	68.849998	5.600000
47	2010-03-01 02:00:00	0.33	0.21	0.84	0.73	0.17	10.810000	14.900000	1.00	74.750000	7.890000
59	2010-03-01 03:00:00	0.38	0.16	0.64	1.00	0.26	2.750000	4.200000	1.00	93.629997	5.130000
...
947	2010-03-02 16:00:00	1.44	0.20	1.86	1.78	0.46	17.090000	20.230000	1.88	91.129997	14.360000
959	2010-03-02 16:00:00	0.64	0.31	1.21	1.78	0.23	28.459999	38.939999	1.88	80.120003	10.740000
971	2010-03-02 17:00:00	0.68	0.18	1.30	1.57	0.43	17.090000	19.990000	1.89	72.089996	13.100000
983	2010-03-02 17:00:00	0.65	0.31	1.39	1.57	0.21	33.500000	45.240002	1.89	73.099998	14.500000
995	2010-03-02 18:00:00	0.51	0.20	0.91	1.27	0.39	20.330000	22.940001	1.42	86.410004	14.280000

83 rows × 17 columns

In [45]:

```
b.columns
```

Out[45]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
      'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],  
      dtype='object')
```


In [46]:

```
b=b[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
      'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
b
```

Out[46]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PXY
11	0.78	0.18	0.84	0.73	0.28	10.420000	11.900000	1.00	90.309998	18.370001	0.81
23	0.70	0.23	1.00	0.73	0.18	17.820000	22.290001	1.00	70.550003	23.639999	0.81
35	0.58	0.17	0.84	0.73	0.28	3.500000	4.950000	1.00	68.849998	5.600000	0.81
47	0.33	0.21	0.84	0.73	0.17	10.810000	14.900000	1.00	74.750000	7.890000	0.81
59	0.38	0.16	0.64	1.00	0.26	2.750000	4.200000	1.00	93.629997	5.130000	0.79
...
947	1.44	0.20	1.86	1.78	0.46	17.090000	20.230000	1.88	91.129997	14.360000	1.84
959	0.64	0.31	1.21	1.78	0.23	28.459999	38.939999	1.88	80.120003	10.740000	1.84
971	0.68	0.18	1.30	1.57	0.43	17.090000	19.990000	1.89	72.089996	13.100000	1.26
983	0.65	0.31	1.39	1.57	0.21	33.500000	45.240002	1.89	73.099998	14.500000	1.26
995	0.51	0.20	0.91	1.27	0.39	20.330000	22.940001	1.42	86.410004	14.280000	0.98

83 rows × 15 columns

In [47]:

```
x=b.iloc[:,0:10]
y=b.iloc[:, -1]
```

In [48]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [49]:

```
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[49]:

LinearRegression()

In [50]:

```
print(lr.intercept_)
```

28079056.293549296

In [51]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[51]:


	Co-efficient
BEN	-52.703785
CO	334.102227
EBE	36.961952
MXY	-7.639622
NMHC	-224.077957
NO_2	-0.909075
NOx	0.088937
OXY	18.311069
O_3	-0.177018
PM10	-0.050966

In [52]:

```
prediction=lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[52]:

<matplotlib.collections.PathCollection at 0x2a95dcdaeb0>



In [53]:

```
print(lr.score(x_test,y_test))
```

0.9083143136566486

In [54]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[54]:

Ridge(alpha=10)

In [55]:

```
rr.score(x_test,y_test)
```

Out[55]:

0.29214401316710437

In [56]:

```
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[56]:

Lasso(alpha=10)

In [57]:

```
la.score(x_test,y_test)
```

Out[57]:

0.26758908780304247

In [58]:

```
from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[58]:

ElasticNet()

In [59]:

```
print(en.coef_)
```

```
[-0.94459074  0.          4.63890773  3.05589677 -2.07547586 -4.83341341
  3.61175272  0.51042608 -0.63410204 -1.68189221]
```

In [60]:

```
print(en.intercept_)
```

28079113.433931526

In [61]:

```
prediction=en.predict(x_test)
print(en.score(x_test,y_test))
```

0.2465375704662155

In [62]:

```
f=StandardScaler().fit_transform(x)
```

In [63]:

```
logr=LogisticRegression()
logr.fit(f,y)
```

Out[63]:

LogisticRegression()

In [64]:

```
g=[[10,20,30,40,50,60,70,80,90,10]]
```

In [65]:

```
prediction=logr.predict(g)
print(prediction)
```

[28079024]

In [66]:

```
logr.classes_
```

Out[66]:

array([28079024, 28079099], dtype=int64)

In [67]:

```
logr.predict_proba(g)[0][0]
```

Out[67]:

1.0

In [68]:

```
logr.score(x_test,y_test)
```

Out[68]:

0.56

In [69]:

```
rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[69]:

```
RandomForestClassifier()
```

In [70]:

```
parameters={'max_depth':[1,2,3,4,5],  
            'min_samples_leaf':[5,10,15,20,25],  
            'n_estimators':[10,20,30,40,50]  
            }
```

In [71]:

```
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

Out[71]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
             param_grid={'max_depth': [1, 2, 3, 4, 5],  
                         'min_samples_leaf': [5, 10, 15, 20, 25],  
                         'n_estimators': [10, 20, 30, 40, 50]},  
             scoring='accuracy')
```

In [72]:

```
grid_search.best_score_
```

Out[72]:

```
1.0
```

In [73]:

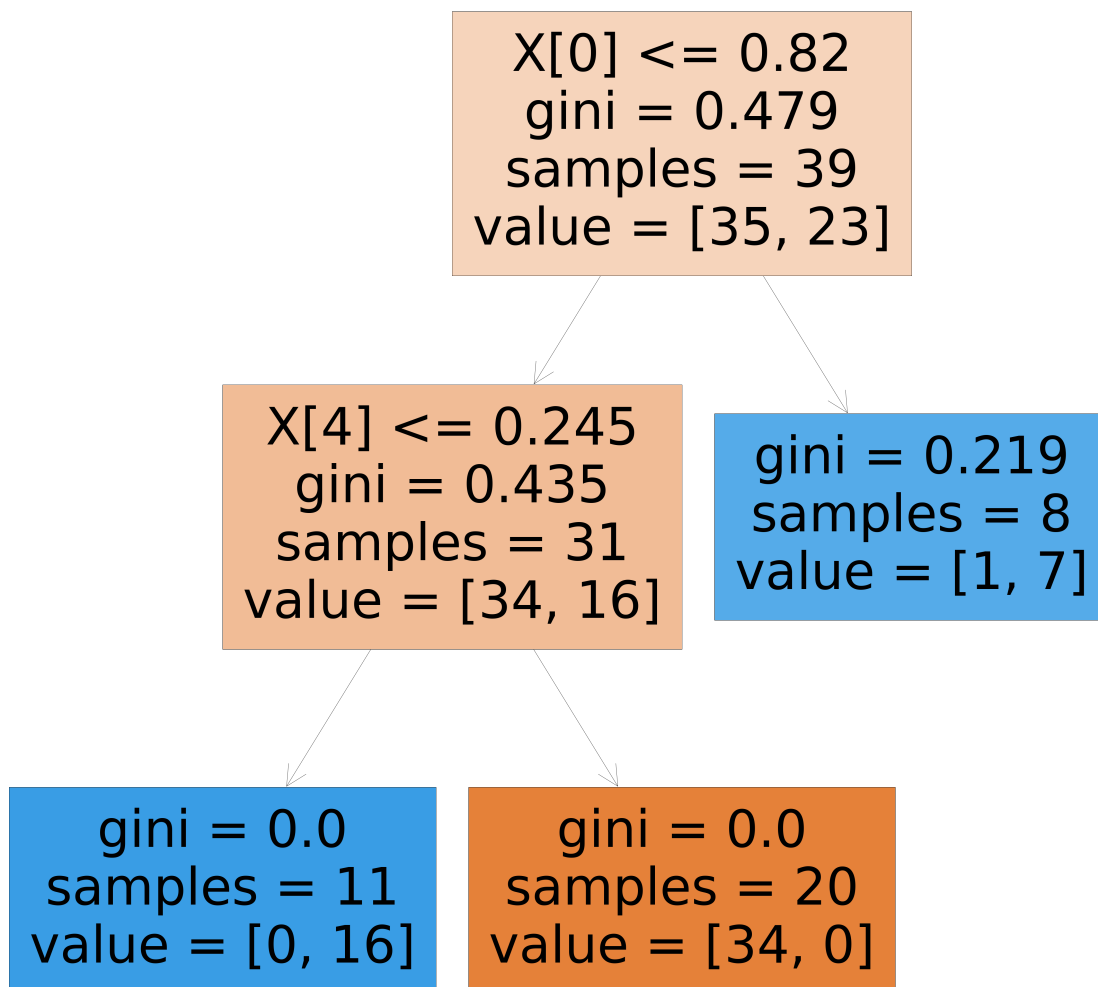
```
rfc_best=grid_search.best_estimator_
```

In [74]:

```
plt.figure(figsize=(80,80))  
plot_tree(rfc_best.estimators_[5],filled=True)
```

Out[74]:

```
[Text(2678.3999999999996, 3624.0, 'X[0] <= 0.82\ngini = 0.479\nsamples = 39\nvalue = [35, 23]'),  
Text(1785.6, 2174.4, 'X[4] <= 0.245\ngini = 0.435\nsamples = 31\nvalue = [34, 16]'),  
Text(892.8, 724.7999999999997, 'gini = 0.0\nsamples = 11\nvalue = [0, 16]'),  
Text(2678.3999999999996, 724.7999999999997, 'gini = 0.0\nsamples = 20\nvalue = [34, 0]'),  
Text(3571.2, 2174.4, 'gini = 0.219\nsamples = 8\nvalue = [1, 7]')]
```



Conclusion: RandomForest score=1.0. This has the highest accuracy.

Madrid 2011

In [75]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\madrid_2011.csv")
a
```

Out[75]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOI
0	2011-11-01 01:00:00	NaN	1.0	NaN	NaN	154.0	84.0	NaN	NaN	NaN	6.0	NaN	NaN
1	2011-11-01 01:00:00	2.5	0.4	3.5	0.26	68.0	92.0	3.0	40.0	24.0	9.0	1.54	8.0
2	2011-11-01 01:00:00	2.9	NaN	3.8	NaN	96.0	99.0	NaN	NaN	NaN	NaN	NaN	7.0
3	2011-11-01 01:00:00	NaN	0.6	NaN	NaN	60.0	83.0	2.0	NaN	NaN	NaN	NaN	NaN
4	2011-11-01 01:00:00	NaN	NaN	NaN	NaN	44.0	62.0	3.0	NaN	NaN	3.0	NaN	NaN
...
209923	2011-09-01 00:00:00	NaN	0.2	NaN	NaN	5.0	19.0	44.0	NaN	NaN	NaN	NaN	NaN
209924	2011-09-01 00:00:00	NaN	0.1	NaN	NaN	6.0	29.0	NaN	11.0	NaN	7.0	NaN	NaN
209925	2011-09-01 00:00:00	NaN	NaN	NaN	0.23	1.0	21.0	28.0	NaN	NaN	NaN	1.44	NaN
209926	2011-09-01 00:00:00	NaN	NaN	NaN	NaN	3.0	15.0	48.0	NaN	NaN	NaN	NaN	NaN
209927	2011-09-01 00:00:00	NaN	NaN	NaN	NaN	4.0	33.0	38.0	13.0	NaN	NaN	NaN	NaN

209928 rows × 14 columns



In [76]:

```
a=a.head(2000)
a
```

Out[76]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL
0	2011-11-01 01:00:00	NaN	1.0	NaN	NaN	154.0	84.0	NaN	NaN	NaN	6.0	NaN	NaN
1	2011-11-01 01:00:00	2.5	0.4	3.5	0.26	68.0	92.0	3.0	40.0	24.0	9.0	1.54	8.7
2	2011-11-01 01:00:00	2.9	NaN	3.8	NaN	96.0	99.0	NaN	NaN	NaN	NaN	NaN	7.2
3	2011-11-01 01:00:00	NaN	0.6	NaN	NaN	60.0	83.0	2.0	NaN	NaN	NaN	NaN	NaN
4	2011-11-01 01:00:00	NaN	NaN	NaN	NaN	44.0	62.0	3.0	NaN	NaN	3.0	NaN	NaN
...
1995	2011-11-04 12:00:00	NaN	0.4	NaN	NaN	23.0	52.0	24.0	NaN	NaN	NaN	NaN	NaN
1996	2011-11-04 12:00:00	NaN	NaN	NaN	NaN	15.0	43.0	20.0	NaN	NaN	3.0	NaN	NaN
1997	2011-11-04 12:00:00	0.2	0.5	0.2	NaN	21.0	47.0	18.0	13.0	NaN	2.0	NaN	1.2
1998	2011-11-04 12:00:00	0.7	0.3	2.0	0.16	10.0	35.0	26.0	10.0	7.0	2.0	1.34	3.1
1999	2011-11-04 12:00:00	NaN	NaN	NaN	0.17	7.0	31.0	43.0	NaN	NaN	NaN	1.28	NaN

2000 rows × 14 columns



In [77]:

```
b=a.dropna()  
b
```

Out[77]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
1	2011-11-01 01:00:00	2.5	0.4	3.5	0.26	68.0	92.0	3.0	40.0	24.0	9.0	1.54	8.7	28
6	2011-11-01 01:00:00	0.7	0.3	1.1	0.16	17.0	66.0	7.0	22.0	16.0	2.0	1.36	1.7	28
25	2011-11-01 02:00:00	1.8	0.3	2.8	0.20	34.0	76.0	3.0	34.0	21.0	8.0	1.71	7.4	28
30	2011-11-01 02:00:00	1.0	0.4	1.3	0.18	31.0	67.0	5.0	25.0	18.0	3.0	1.40	2.9	28
49	2011-11-01 03:00:00	1.3	0.2	2.4	0.22	29.0	72.0	3.0	33.0	20.0	8.0	1.75	6.2	28
...
1950	2011-11-04 10:00:00	1.0	0.4	2.2	0.22	31.0	61.0	10.0	16.0	13.0	3.0	1.78	2.9	28
1969	2011-11-04 11:00:00	1.9	0.3	2.5	0.15	38.0	57.0	16.0	21.0	11.0	8.0	1.31	6.6	28
1974	2011-11-04 11:00:00	0.7	0.3	2.5	0.20	15.0	43.0	19.0	10.0	7.0	2.0	1.76	3.3	28
1993	2011-11-04 12:00:00	1.0	0.3	1.6	0.14	30.0	55.0	24.0	18.0	10.0	7.0	1.40	4.9	28
1998	2011-11-04 12:00:00	0.7	0.3	2.0	0.16	10.0	35.0	26.0	10.0	7.0	2.0	1.34	3.1	28

165 rows × 14 columns

In [78]:

```
b.columns
```

Out[78]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
      'SO_2', 'TCH', 'TOL', 'station'],  
      dtype='object')
```

In [79]:

```
b=b[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
      'SO_2', 'TCH', 'TOL', 'station']]
b
```

Out[79]:

	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
1	2.5	0.4	3.5	0.26	68.0	92.0	3.0	40.0	24.0	9.0	1.54	8.7	28079008
6	0.7	0.3	1.1	0.16	17.0	66.0	7.0	22.0	16.0	2.0	1.36	1.7	28079024
25	1.8	0.3	2.8	0.20	34.0	76.0	3.0	34.0	21.0	8.0	1.71	7.4	28079008
30	1.0	0.4	1.3	0.18	31.0	67.0	5.0	25.0	18.0	3.0	1.40	2.9	28079024
49	1.3	0.2	2.4	0.22	29.0	72.0	3.0	33.0	20.0	8.0	1.75	6.2	28079008
...
1950	1.0	0.4	2.2	0.22	31.0	61.0	10.0	16.0	13.0	3.0	1.78	2.9	28079024
1969	1.9	0.3	2.5	0.15	38.0	57.0	16.0	21.0	11.0	8.0	1.31	6.6	28079008
1974	0.7	0.3	2.5	0.20	15.0	43.0	19.0	10.0	7.0	2.0	1.76	3.3	28079024
1993	1.0	0.3	1.6	0.14	30.0	55.0	24.0	18.0	10.0	7.0	1.40	4.9	28079008
1998	0.7	0.3	2.0	0.16	10.0	35.0	26.0	10.0	7.0	2.0	1.34	3.1	28079024

165 rows × 13 columns

In [80]:

```
x=b.iloc[:,0:5]
y=b.iloc[:, -1]
```

In [81]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [82]:

```
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[82]:

LinearRegression()

In [83]:

```
print(lr.intercept_)
```

28079004.011541203

In [84]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[84]:


	Co-efficient
BEN	-8.293828
CO	97.736483
EBE	2.596222
NMHC	7.584429
NO	-0.576532

In [85]:

```
prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[85]:

<matplotlib.collections.PathCollection at 0x2a95e29cee0>



In [86]:

```
print(lr.score(x_test,y_test))
```

0.2983438088128415

In [87]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[87]:

Ridge(alpha=10)

In [88]:

```
rr.score(x_test,y_test)
```

Out[88]:

0.10949990447518754

In [89]:

```
la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

Out[89]:

Lasso(alpha=10)

In [90]:

```
la.score(x_test,y_test)
```

Out[90]:

0.1342859246108833

In [91]:

```
en=ElasticNet()  
en.fit(x_train,y_train)
```

Out[91]:

ElasticNet()

In [92]:

```
print(en.coef_)
```

[0. 0. 0.24960629 0. -0.24491266]

In [93]:

```
print(en.intercept_)
```

28079019.37398652

In [94]:

```
prediction=en.predict(x_test)  
print(en.score(x_test,y_test))
```

0.11137093683449806

In [95]:

```
f=StandardScaler().fit_transform(x)
```

In [96]:

```
logr=LogisticRegression()  
logr.fit(f,y)
```

Out[96]:

```
LogisticRegression()
```

In [97]:

```
g=[[10,20,30,40,50]]
```

In [98]:

```
prediction=logr.predict(g)  
print(prediction)
```

```
[28079008]
```

In [99]:

```
logr.classes_
```

Out[99]:

```
array([28079008, 28079024], dtype=int64)
```

In [100]:

```
logr.predict_proba(g)[0][0]
```

Out[100]:

```
1.0
```

In [101]:

```
logr.score(x_test,y_test)
```

Out[101]:

```
0.5
```

In [102]:

```
rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[102]:

```
RandomForestClassifier()
```

In [103]:

```
parameters={'max_depth':[1,2,3,4,5],  
            'min_samples_leaf':[5,10,15,20,25],  
            'n_estimators':[10,20,30,40,50]  
            }
```

In [104]:

```
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

Out[104]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),
             param_grid={'max_depth': [1, 2, 3, 4, 5],
                          'min_samples_leaf': [5, 10, 15, 20, 25],
                          'n_estimators': [10, 20, 30, 40, 50]},
             scoring='accuracy')
```

In [105]:

```
grid_search.best_score_
```

Out[105]:

```
0.8867211131276467
```

In [106]:

```
rfc_best=grid_search.best_estimator_
```

In [107]:

```
plt.figure(figsize=(80,80))
plot_tree(rfc_best.estimators_[5],filled=True)
```

Out[107]:

```
[Text(2637.818181818182, 3913.92, 'X[2] <= 1.45\ngini = 0.5\nsamples = 68\nvalue = [57, 58]'),
 Text(1623.2727272727273, 3044.1600000000003, 'X[4] <= 2.5\ngini = 0.48\nsamples = 55\nvalue = [36, 54]'),
 Text(811.6363636363636, 2174.4, 'X[0] <= 0.45\ngini = 0.048\nsamples = 25\nvalue = [1, 40]'),
 Text(405.8181818181818, 1304.6400000000003, 'gini = 0.165\nsamples = 7\nvalue = [1, 10]'),
 Text(1217.4545454545455, 1304.6400000000003, 'gini = 0.0\nsamples = 18\nvalue = [0, 30]'),
 Text(2434.909090909091, 2174.4, 'X[3] <= 0.185\ngini = 0.408\nsamples = 30\nvalue = [35, 14]'),
 Text(2029.090909090909, 1304.6400000000003, 'X[4] <= 7.5\ngini = 0.272\nsamples = 24\nvalue = [31, 6]'),
 Text(1623.2727272727273, 434.8800000000001, 'gini = 0.496\nsamples = 9\nvalue = [6, 5]'),
 Text(2434.909090909091, 434.8800000000001, 'gini = 0.074\nsamples = 15\nvalue = [25, 1]'),
 Text(2840.7272727272725, 1304.6400000000003, 'gini = 0.444\nsamples = 6\nvalue = [4, 8]'),
 Text(3652.3636363636365, 3044.1600000000003, 'X[1] <= 0.25\ngini = 0.269\nsamples = 13\nvalue = [21, 4]'),
 Text(3246.5454545454545, 2174.4, 'gini = 0.0\nsamples = 5\nvalue = [8, 0]'),
 Text(4058.181818181818, 2174.4, 'gini = 0.36\nsamples = 8\nvalue = [13, 4]')]
```

Conclusion: RandomForest Score=0.8870235934664247. It has the highest accuracy.

Madrid 2012

In [108]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\C10-air\madrid_2012.csv")
a
```

Out[108]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL
	2012-09-01 01:00:00	0.3	0.3	0.7	NaN	3.0	18.0	55.0	10.0	9.0	1.0	NaN	2.4
1	2012-09-01 01:00:00	0.3	0.3	0.7	NaN	3.0	18.0	55.0	10.0	9.0	1.0	NaN	2.4
2	2012-09-01 01:00:00	0.4	NaN	NaN	NaN	2.0	18.0	55.0	10.0	9.0	1.0	NaN	2.4
3	2012-09-01 01:00:00	NaN	0.2	NaN	NaN	1.0	6.0	50.0	NaN	NaN	NaN	NaN	NaN
4	2012-09-01 01:00:00	NaN	NaN	NaN	NaN	1.0	13.0	54.0	NaN	NaN	3.0	NaN	NaN
...
210715	2012-03-01 00:00:00	NaN	0.6	NaN	NaN	37.0	84.0	14.0	NaN	NaN	NaN	NaN	NaN
210716	2012-03-01 00:00:00	NaN	0.4	NaN	NaN	5.0	76.0	NaN	17.0	NaN	7.0	NaN	NaN
210717	2012-03-01 00:00:00	NaN	NaN	NaN	0.34	3.0	41.0	24.0	NaN	NaN	NaN	1.34	NaN
210718	2012-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	44.0	36.0	NaN	NaN	NaN	NaN	NaN
210719	2012-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	56.0	40.0	18.0	NaN	NaN	NaN	NaN

210720 rows × 14 columns

In [109]:

```
a=a.head(2000)
a
```

Out[109]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2012-09-01 01:00:00	NaN	0.2	NaN	NaN	7.0	18.0	NaN	NaN	NaN	2.0	NaN	NaN	2i
1	2012-09-01 01:00:00	0.3	0.3	0.7	NaN	3.0	18.0	55.0	10.0	9.0	1.0	NaN	2.4	2i
2	2012-09-01 01:00:00	0.4	NaN	0.7	NaN	2.0	10.0	NaN	NaN	NaN	NaN	NaN	1.5	2
3	2012-09-01 01:00:00	NaN	0.2	NaN	NaN	1.0	6.0	50.0	NaN	NaN	NaN	NaN	NaN	2i
4	2012-09-01 01:00:00	NaN	NaN	NaN	NaN	1.0	13.0	54.0	NaN	NaN	3.0	NaN	NaN	2i
...	
1995	2012-09-04 12:00:00	NaN	0.4	NaN	NaN	3.0	10.0	59.0	NaN	NaN	NaN	NaN	NaN	2i
1996	2012-09-04 12:00:00	NaN	NaN	NaN	NaN	4.0	15.0	61.0	NaN	NaN	3.0	NaN	NaN	2i
1997	2012-09-04 12:00:00	0.1	0.3	1.0	NaN	7.0	19.0	57.0	20.0	NaN	1.0	NaN	0.7	2i
1998	2012-09-04 12:00:00	0.5	0.2	1.2	0.24	3.0	12.0	63.0	13.0	8.0	3.0	1.35	1.2	2i
1999	2012-09-04 12:00:00	NaN	NaN	NaN	0.11	1.0	6.0	78.0	NaN	NaN	NaN	1.18	NaN	2i

2000 rows × 14 columns



In [110]:

```
b=a.dropna()  
b
```

Out[110]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	station
6	2012-09-01 01:00:00	0.4	0.2	0.8	0.24	1.0	7.0	57.0	11.0	7.0	2.0	1.33	0.6	280
30	2012-09-01 02:00:00	0.4	0.2	0.7	0.24	1.0	5.0	55.0	5.0	5.0	2.0	1.33	0.5	280
54	2012-09-01 03:00:00	0.4	0.2	0.7	0.24	1.0	4.0	56.0	6.0	4.0	2.0	1.33	0.5	280
78	2012-09-01 04:00:00	0.3	0.2	0.7	0.25	1.0	5.0	54.0	6.0	5.0	2.0	1.34	0.4	280
102	2012-09-01 05:00:00	0.4	0.2	0.7	0.24	1.0	3.0	53.0	8.0	5.0	2.0	1.33	0.5	280
...
1902	2012-09-04 08:00:00	0.4	0.2	0.8	0.23	1.0	13.0	51.0	7.0	8.0	2.0	1.34	0.4	280
1926	2012-09-04 09:00:00	0.6	0.2	0.9	0.24	2.0	21.0	46.0	12.0	8.0	2.0	1.35	0.6	280
1950	2012-09-04 10:00:00	0.6	0.2	1.0	0.24	4.0	19.0	47.0	13.0	9.0	2.0	1.35	0.9	280
1974	2012-09-04 11:00:00	0.5	0.2	1.2	0.25	4.0	15.0	54.0	15.0	9.0	2.0	1.36	1.2	280
1998	2012-09-04 12:00:00	0.5	0.2	1.2	0.24	3.0	12.0	63.0	13.0	8.0	3.0	1.35	1.2	280

84 rows × 14 columns

In [111]:

```
b.columns
```

Out[111]:

```
Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
      'SO_2', 'TCH', 'TOL', 'station'],  
      dtype='object')
```

In [112]:

```
c=b[['BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',  
    'SO_2', 'TCH', 'TOL']]  
d=b['station']
```

In [117]:

```
x=c.iloc[:,0:11]  
y=c.iloc[:, -1]
```

In [118]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [119]:

```
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[119]:

LinearRegression()

In [120]:

```
print(lr.score(x_test,y_test))
```

0.698411509777724

In [121]:

```
rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[121]:

Ridge(alpha=10)

In [122]:

```
rr.score(x_test,y_test)
```

Out[122]:

0.44872555626777366

In [123]:

```
la=Lasso()  
la.fit(x_train,y_train)
```

Out[123]:

Lasso()

In [124]:

```
la.score(x_test,y_test)
```

Out[124]:

-0.150303495066209

In [125]:

```
en=ElasticNet()  
en.fit(x_train,y_train)
```

Out[125]:

ElasticNet()

In [126]:

```
prediction=en.predict(x_test)  
print(en.score(x_test,y_test))
```

-0.14983006866203818

In [139]:

```
x1=c.drop('NO',axis=1)  
y1=c['NO']
```

In [170]:

```
x1_train,x1_test,y1_train,y1_test=train_test_split(x1,y1,test_size=0.4,random_state=40)
```

In [171]:

```
f=StandardScaler().fit_transform(x1)
```

In [172]:

```
logr=LogisticRegression(max_iter=10000)  
logr.fit(x1,y1)
```

Out[172]:

LogisticRegression(max_iter=10000)

In [173]:

```
g=[[10,20,30,40,50,60,70,80,90,100,110]]
```

In [174]:

```
prediction=logr.predict(g)  
print(prediction)
```

[3.]

In [175]:

```
logr.predict_proba(g)[0][0]
```

Out[175]:

```
1.1971510187584438e-121
```

In [177]:

```
logr.score(x1_test,y1_test)
```

Out[177]:

```
0.9117647058823529
```

In [178]:

```
rfc=RandomForestClassifier()  
rfc.fit(x1_train,y1_train)
```

Out[178]:

```
RandomForestClassifier()
```

In [179]:

```
parameters={'max_depth':[1,2,3,4,5,6],  
            'min_samples_leaf':[5,10,15,20,25],  
            'n_estimators':[10,20,30,40,50]  
            }
```

In [180]:

```
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x1_train,y1_train)
```

Out[180]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
             param_grid={'max_depth': [1, 2, 3, 4, 5, 6],  
                         'min_samples_leaf': [5, 10, 15, 20, 25],  
                         'n_estimators': [10, 20, 30, 40, 50]},  
             scoring='accuracy')
```

In [181]:

```
grid_search.best_score_
```

Out[181]:

```
0.78
```

In [182]:


```
rfc_best=grid_search.best_estimator_
```

In [183]:

```
plt.figure(figsize=(80,10))  
plot_tree(rfc_best.estimators_[5],filled=True)
```

Out[183]:

```
[Text(2232.0, 407.70000000000005, 'X[4] <= 6.5\nngini = 0.374\nsamples = 35\n\nvalue = [39, 5, 4, 2]'),  
Text(1116.0, 135.89999999999998, 'gini = 0.0\nsamples = 16\n\nvalue = [22, 0, 0, 0]'),  
Text(3348.0, 135.89999999999998, 'gini = 0.574\nsamples = 19\n\nvalue = [17, 5, 4, 2]')]
```



Conclusion: Logistic Score=0.9117647058823529. It has the highest accuracy.

In []: